

Azure Machine Learning Service

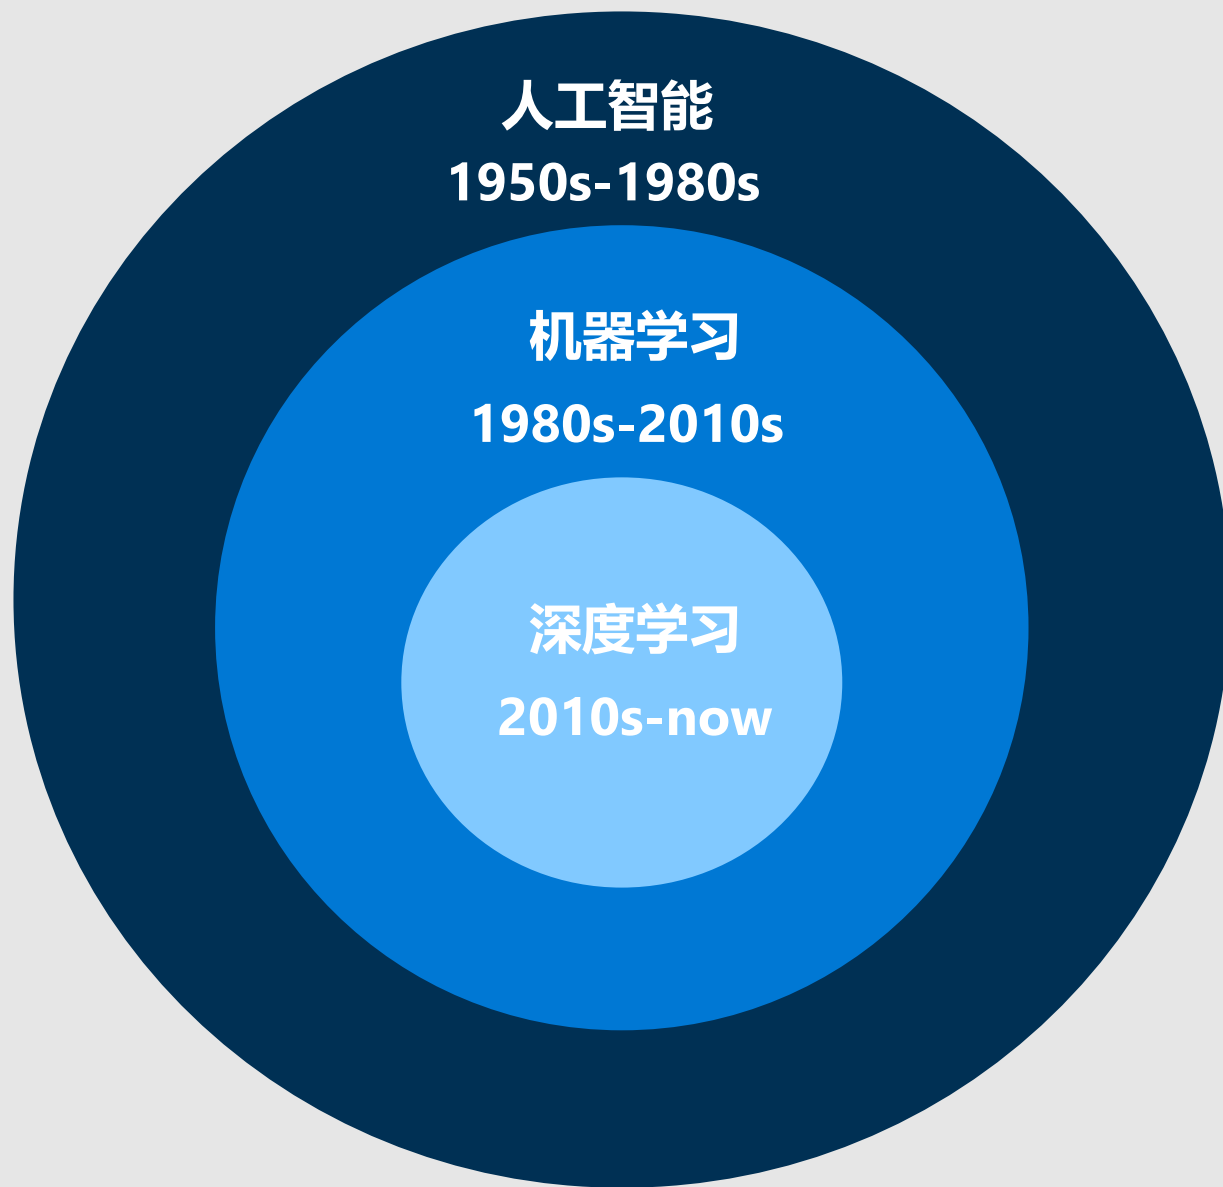
Kinfey.Lo

Microsoft MVP/Xamarin MVP


Github: <https://github.com/lokinfey>

Email: lokinfey@outlook.com

AI 技术不断变化



人工智能的热点方向



计算机视觉



翻译




搜索

如何做AI？



Local PC / Laptop



GPU



Cloud Computing

人工智能技术栈



语言



框架



硬件

放弃AI!!!




Azure 提供的AI能力

AI 应用和服务




Azure Bot Service
Azure Cognitive Services

机器学习



Azure Databricks
Azure Machine Learning

知识挖掘



Azure Cognitive Search

Machine Learning on Azure

提供一些预定义的模型

能降低开发门槛



Vision



Speech



Language



Search

兼容不同的开发工具

快速完成模型开发，简化开发流程



PyCharm



Jupyter



Visual Studio Code



Command line

对人工智能开发框架的支持

根据你的需要创建深度学习的解决方案



Pytorch



TensorFlow



Scikit-Learn



Onnx

提供生产服务

为开发团队提供数据，和训练环境



Azure
Databricks



Azure Machine
Learning



Machine
Learning VMs

强大的硬件架构支持

加速深度学习环境



CPU



GPU



FPGA



Azure Cognitive Services

1. Object Detection & Recognition for thousands of objects
2. Video Indexer (Preview)
3. Speech Recognition with customization
4. Speech Synthesis with customizable voice
5. Speech to Speech Translation
6. Text analytics with entity detection
7. Language Understanding (LUIS) with new integrated offer
8. QnA Maker is Generally Available
9. Bing Visual Search with smart identification
10. Bing Search SDK is Generally Available



Vision




Speech



Language



Conversation



Bing Search



Knowledge

Azure Cognitive Services

Use AI to solve business problems



Vision

Image-processing algorithms to smartly identify, caption, index, and moderate your pictures and videos.



Knowledge

Map complex information and data in order to solve tasks such as intelligent recommendations and semantic search.



Language

Allow your apps to process natural language with pre-built scripts, evaluate sentiment and learn how to recognize what users want.



Speech

Convert spoken audio into text, use voice for verification, or add speaker recognition to your app.



Search

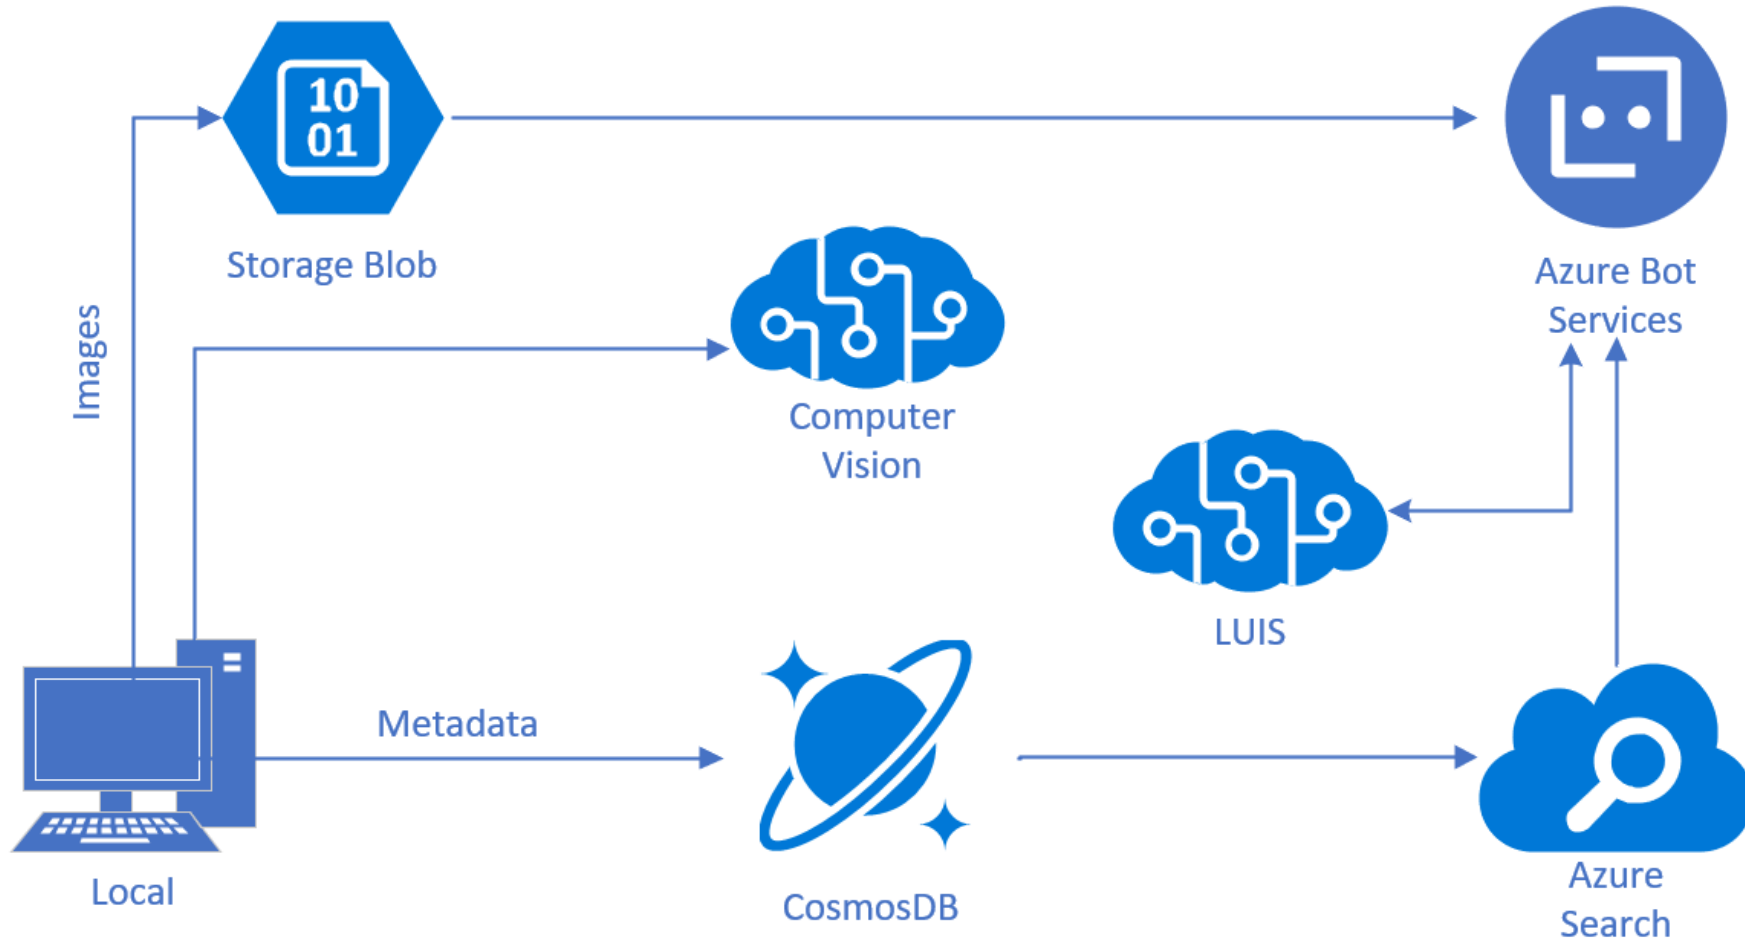
Add Bing Search APIs to your apps and harness the ability to comb billions of webpages, images, videos, and news with a single API call.



Anomaly Detection

Add anomaly detection capabilities to your apps to identify problems as soon as they occur.

Azure Cognitive Services



Azure Cognitive Services

The screenshot displays the Microsoft Azure portal interface. At the top, the header includes the 'Microsoft Azure' logo, a search bar for resources, services, and docs, and user information for 'lokinfey@outlook.com' (LO KINFY).

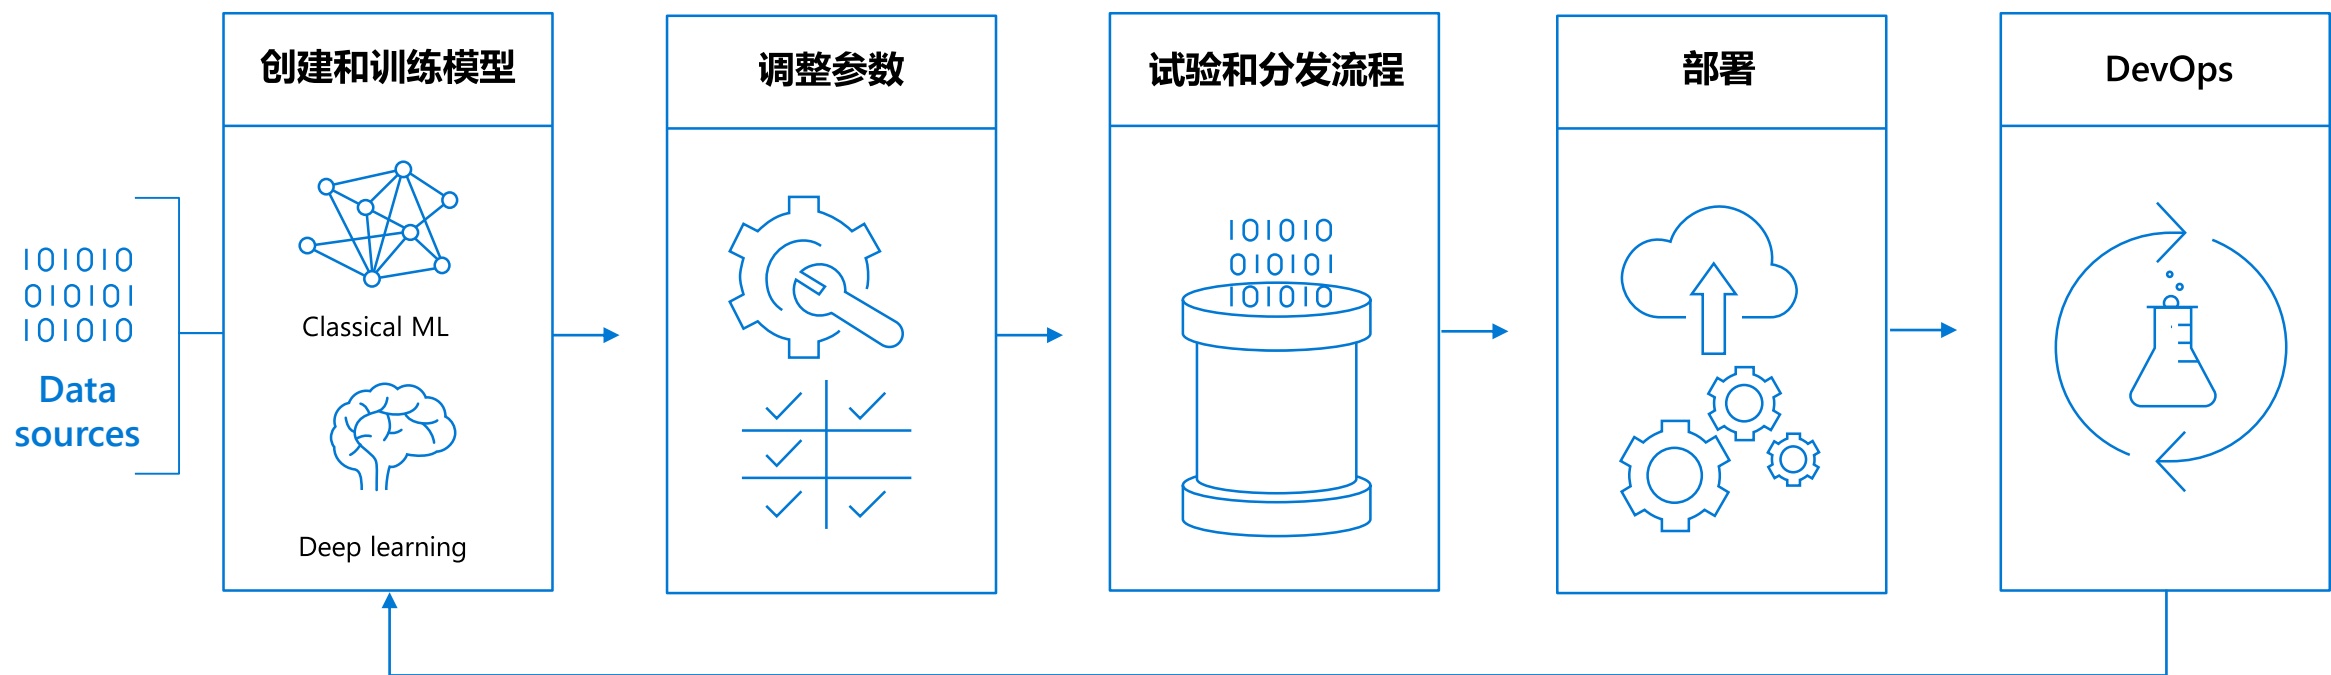
The left sidebar contains navigation links: 'Create a resource', 'Home', 'Dashboard', 'All services', 'FAVORITES', 'Resource groups', 'All resources', 'Recent', 'App Services', 'Virtual machines', 'Virtual machines (classic)', 'SQL databases', 'Cloud services (classic)', 'Subscriptions', 'Azure Active Directory', 'Monitor', 'Security Center', 'Cost Management + Billing', 'Help + support', 'Advisor', and 'Media services'.

The main content area is titled 'New' and features a search bar labeled 'Search the Marketplace'. Below this, there are two columns: 'Azure Marketplace' and 'Featured'. The 'Azure Marketplace' column lists various service categories, with 'AI + Machine Learning' highlighted. The 'Featured' column displays several cognitive services with their respective icons and quickstart tutorials:

- Machine Learning service workspace**: Learn more
- Data Science Virtual Machine - Windows 2016 (preview)**: Quickstart tutorial (marked as PREVIEW)
- Web App Bot**: Quickstart tutorial
- Computer Vision**: Quickstart tutorial
- Face**: Quickstart tutorial
- Text Analytics**: Quickstart tutorial
- Language Understanding**: Quickstart tutorial
- Bing Search v7**: Quickstart tutorial
- Ubuntu Server 18.04 LTS**: Learn more

**可是老板
不会因为用Azure Cognitive Service 感到满足**

构建AI项目的步骤






Azure Machine Learning service

是一个提升生产力的AI开发平台



Boost your data science productivity



Increase your rate of experimentation



Deploy and manage your models everywhere



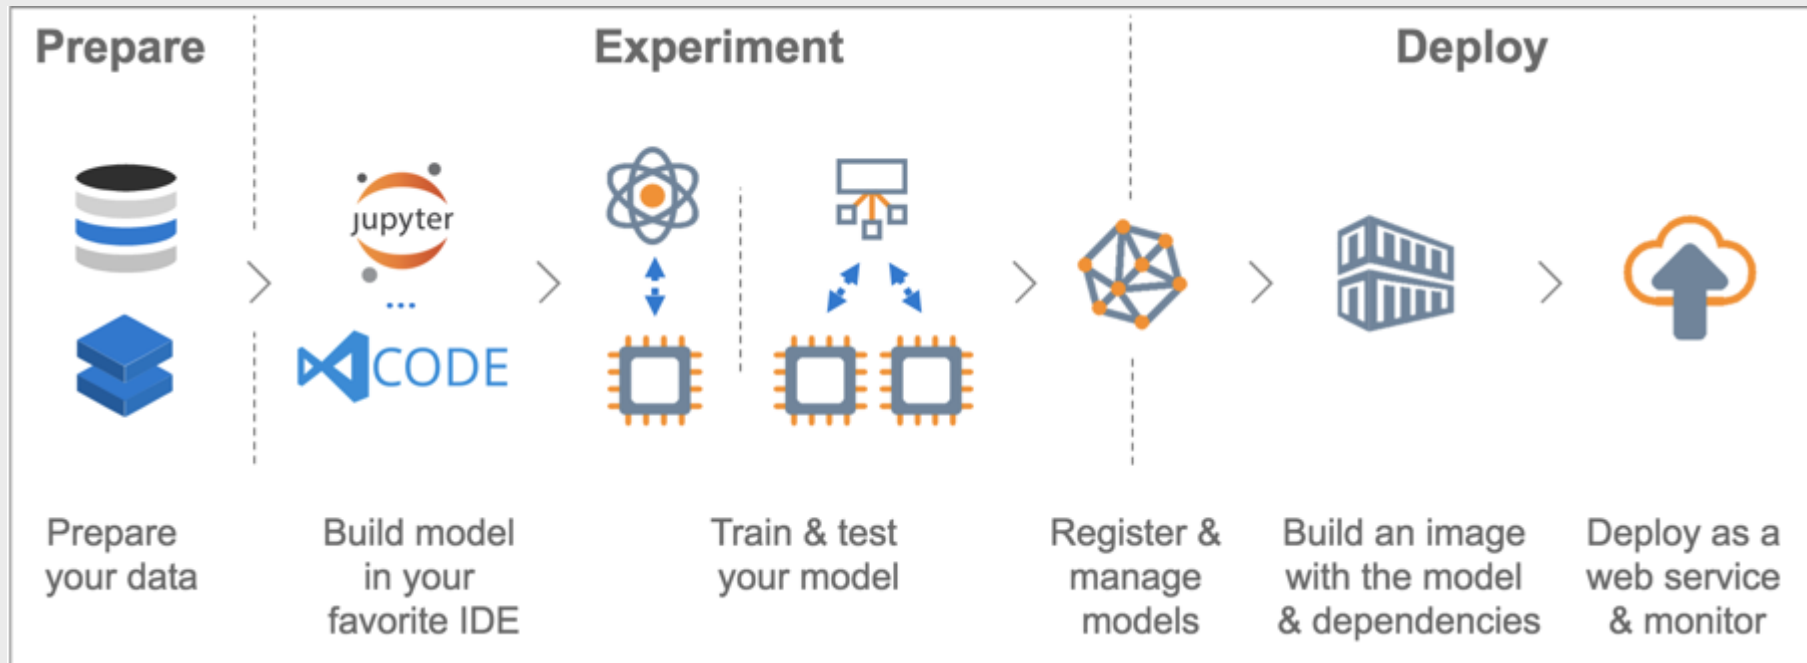
Built with your needs in mind

- Automated machine learning
- Managed compute
- Simple deployment
- DevOps for machine learning
- Support for open source frameworks
- Tool agnostic Python SDK

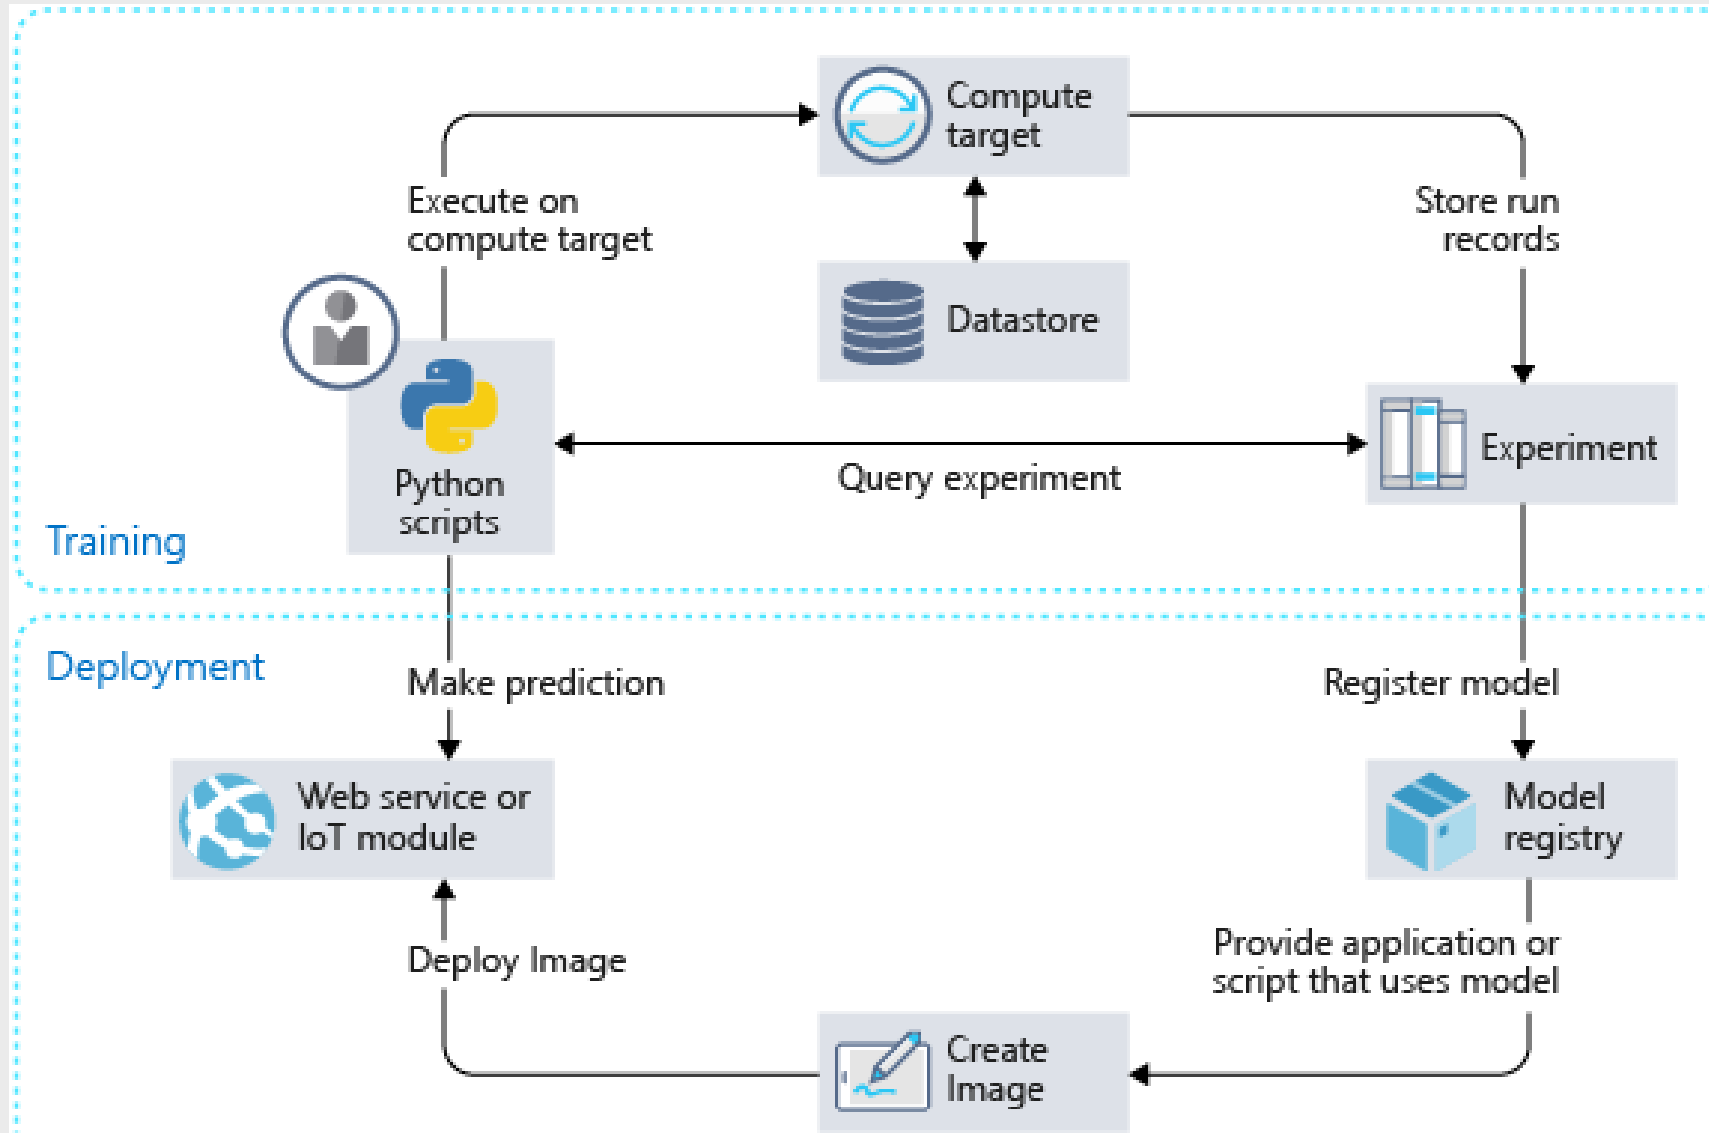
Seamlessly integrated with the Azure Portfolio


Azure Machine Learning Service

Azure Machine Learning service provides a cloud-based environment you can use to prep data, train, test, deploy, manage, and track machine learning models.



Azure Machine Learning Service






Automated machine learning

Azure Machine Learning

Automated machine learning



估算一台车的价格?

模型创建相当费时

Which features?

- Mileage
- Condition
- Car brand
- Year of make
- Regulations
- ...

Which algorithm?

- Gradient Boosted
- Nearest Neighbors
- SVM
- Bayesian Regression
- LGBM
- ...

Which parameters?

- Parameter 1
- Parameter 2
- Random Forest Split
- Random Forest Leaf
- Others

30%
Model

模型创建相当费时

Which features?


- Mileage
- Condition
- Car brand
- Year of make
- Regulations
- ...

Which algorithm?

- Gradient Boosted
- Nearest Neighbors
- SGD
- Bayesian Regression
- LGBM
- ...

Which parameters?

- Neighbors
- Weights
- Min Samples Split
- Min Samples Leaf
- XYZ



↑

Iterate

模型创建相当费时

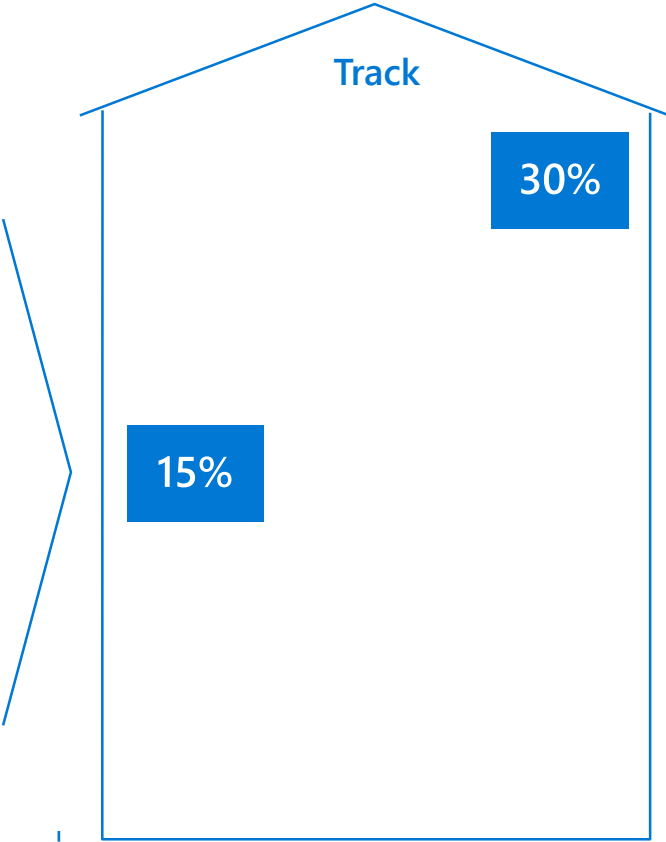

Which features?



Which algorithm?



Which parameters?



Iterate

Automated Machine Learning 加速模型训练

Input

101010
010101
101010

Enter data




Define goals




Apply constraints



Intelligently test multiple models in parallel




Output



Understand the inner workings of ML by analyzing feature importance



不是单一的迭代创建模型，而且可以针对每次训练的模型进行分析，让你更快速找到最佳模型




Managed compute

训练时的基础设施



Dependencies and Containers

Leverage system-managed AML compute or bring your own compute



Distribute data

Manage and share resources across a workspace




Schedule jobs

Train at cloud scale using a framework of choice



Scale resources

Autoscale resources to only pay while running a job



Provision clusters

Use the latest NDv2 series VMs with the NVIDIA V100 GPUs

Powerful infrastructure


Accelerate deep learning



CPUs

General purpose machine learning


D, F, L, M, H Series



GPUs

Deep learning

N Series



FPGAs

Specialized hardware accelerated deep learning

Project Brainwave


Optimized for flexibility

Optimized for performance



FPGA NEW UPDATES:


Support for image classification and recognition scenarios
ResNet 50, ResNet 152, VGG-16, SSD-VGG, DenseNet-121



Azure Machine Learning experiments

DevOps loop for data science

Prepare





Prepare
Data

Experiment



Build model
(your favorite IDE)



Train &
Test Model



Register and
Manage Model

Deploy




Build
Image



Deploy Service
Monitor Model


Experimentation




Leverage service-side capture of run metrics, output logs and models



Manage training jobs locally, scaled-up or scaled-out

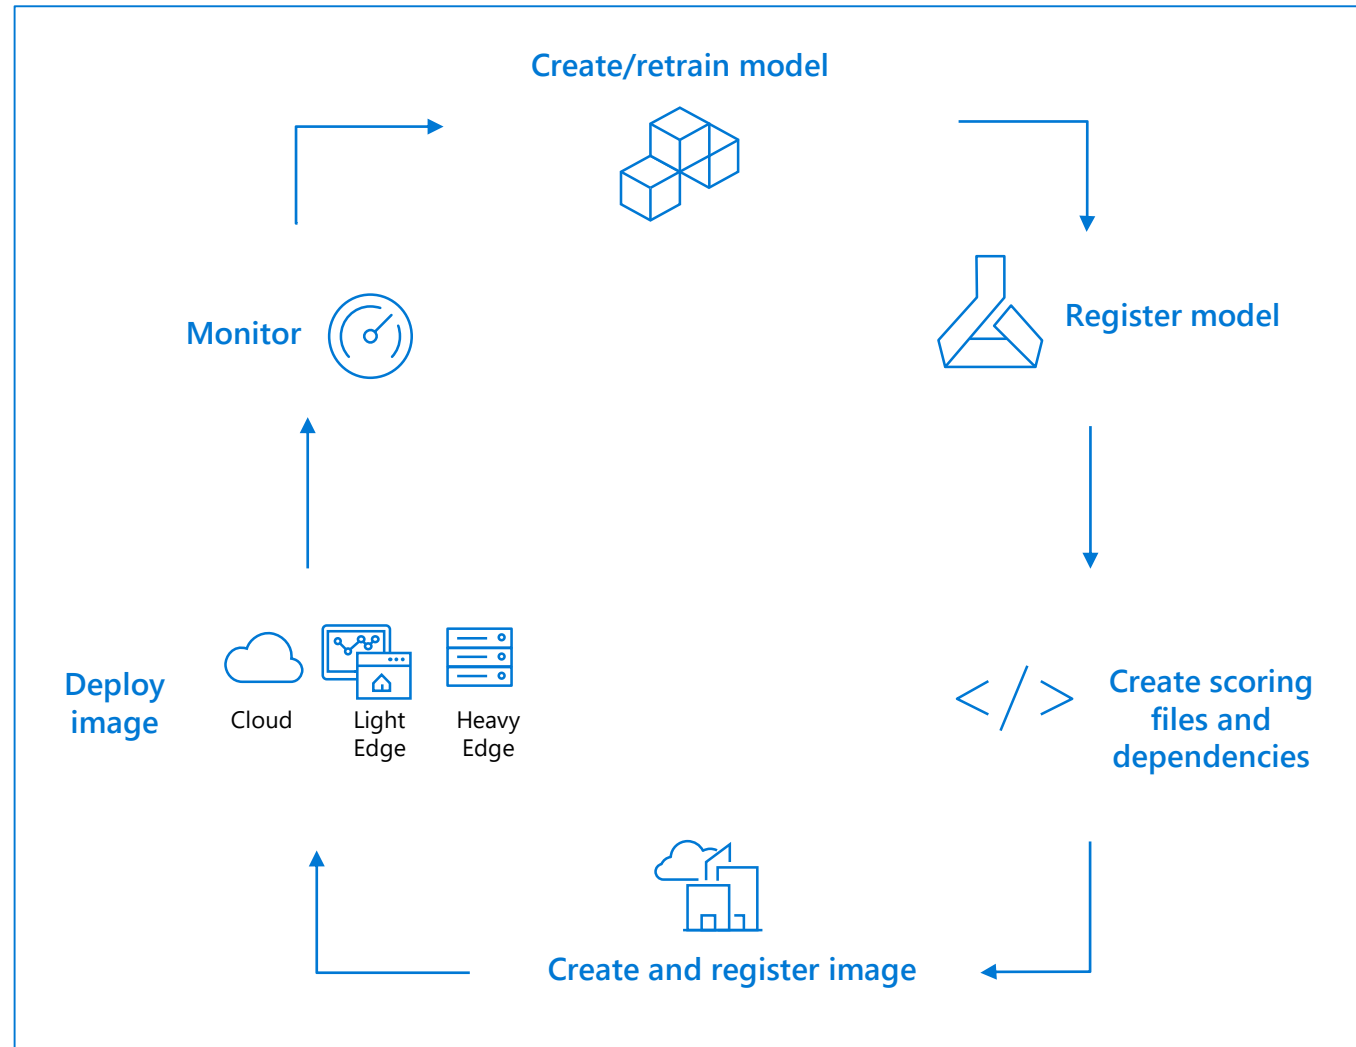


Use leaderboards, side by side run comparison and model selection




Conduct a hyperparameter search on traditional ML or DNN

Model management in Azure Machine Learning




Model management in detail




Create/Retrain Model

Enable DevOps with full CI/CD integration with VSTS



Register Model


Track model versions with a central model registry



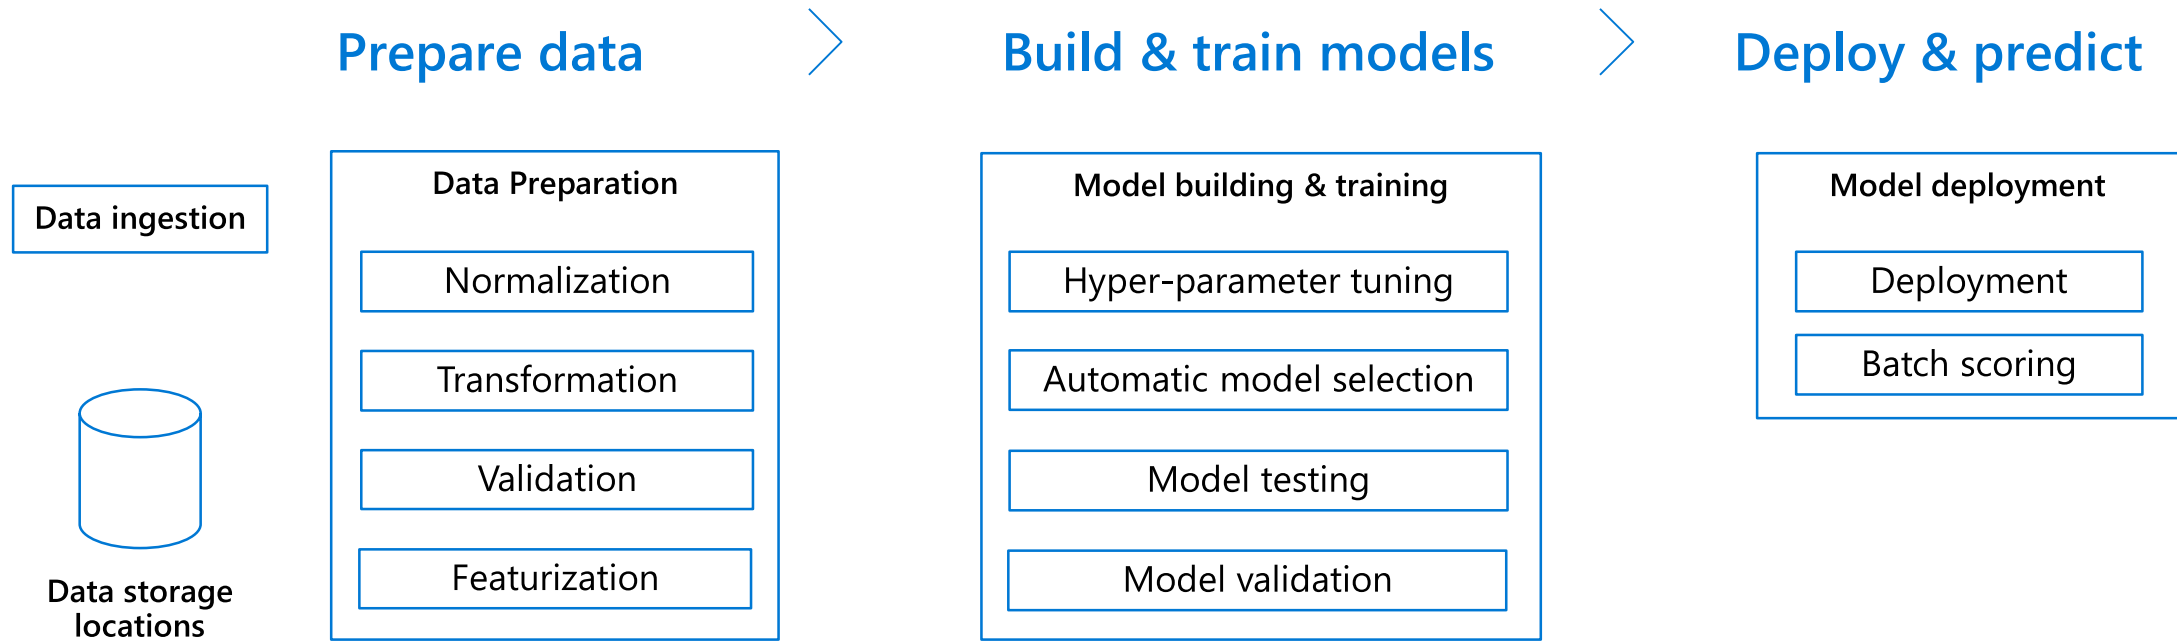
Monitor

Oversea deployments through Azure AppInsights

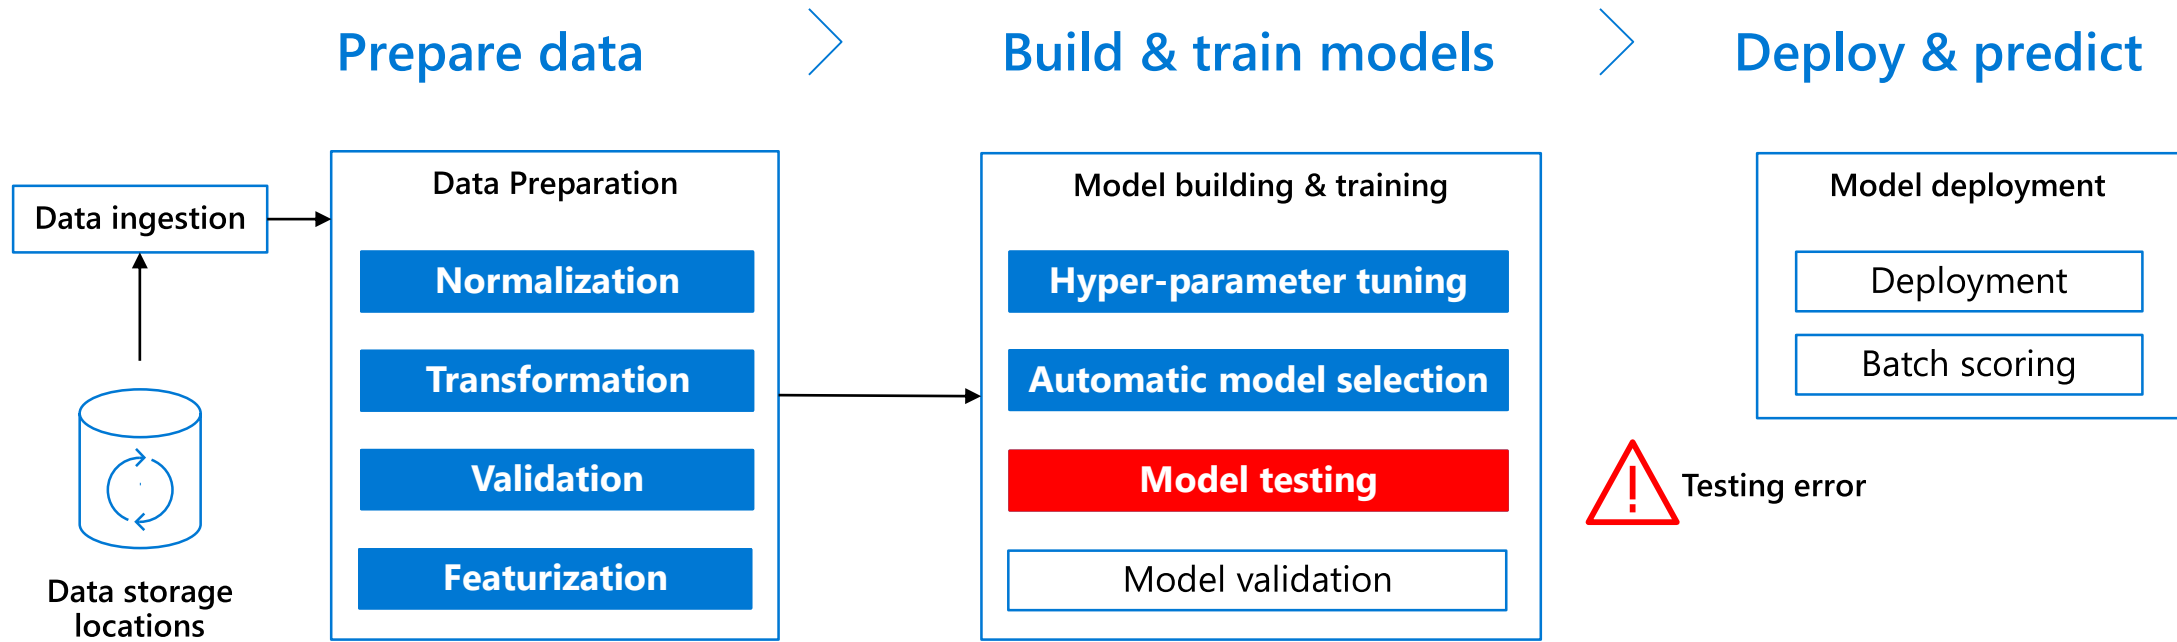
Azure Machine Learning pipelines



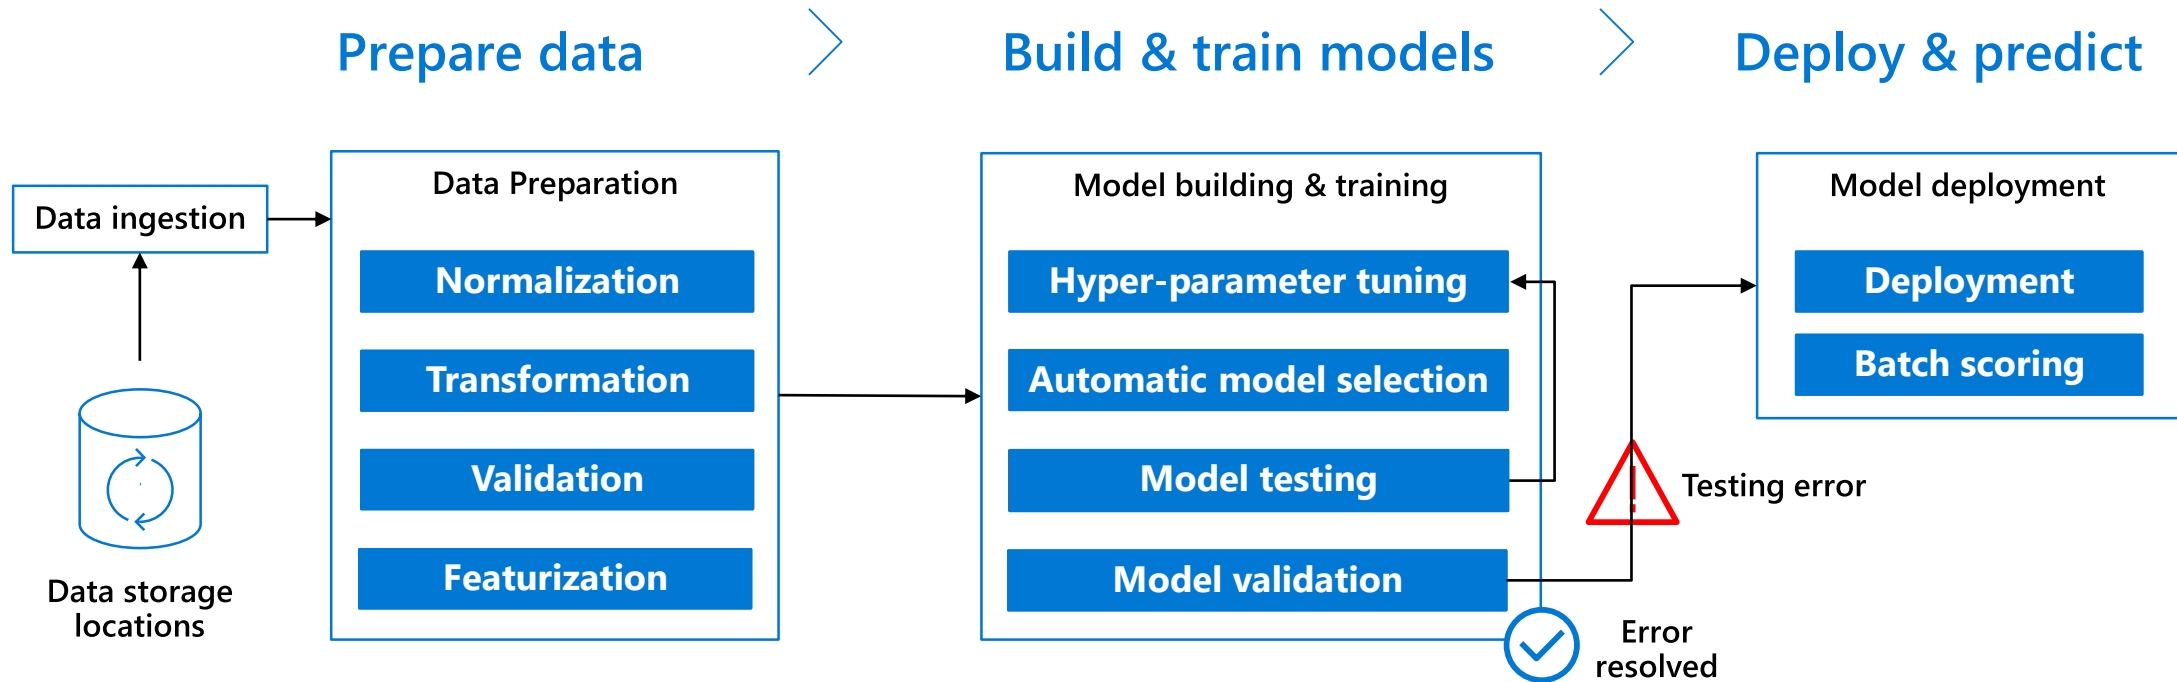
Azure Machine Learning pipelines



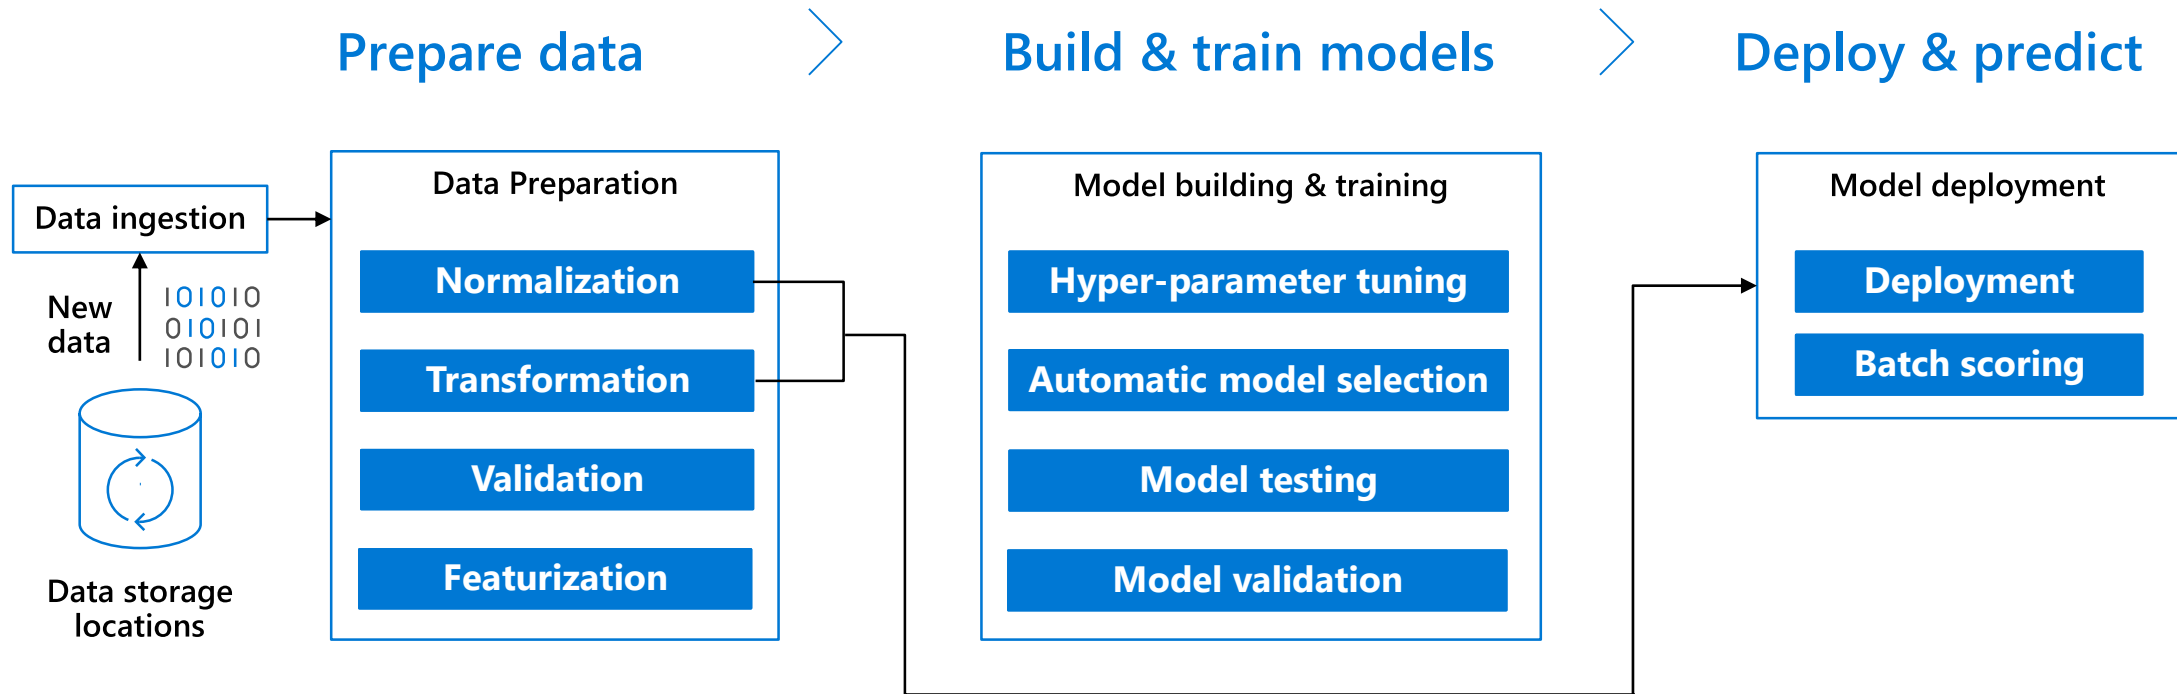
Azure Machine Learning pipelines



Azure Machine Learning pipelines



Azure Machine Learning pipelines with new data




Advantages of Azure ML Pipelines



Unattended runs

Schedule a few steps to run in parallel or in sequence to focus on other tasks while your pipeline runs




Tracking and versioning

Name and version your data sources, inputs and outputs with the pipelines SDK



Reusability

Create templates of pipelines for specific scenarios such as retraining and batch scoring



Mixed and diverse compute

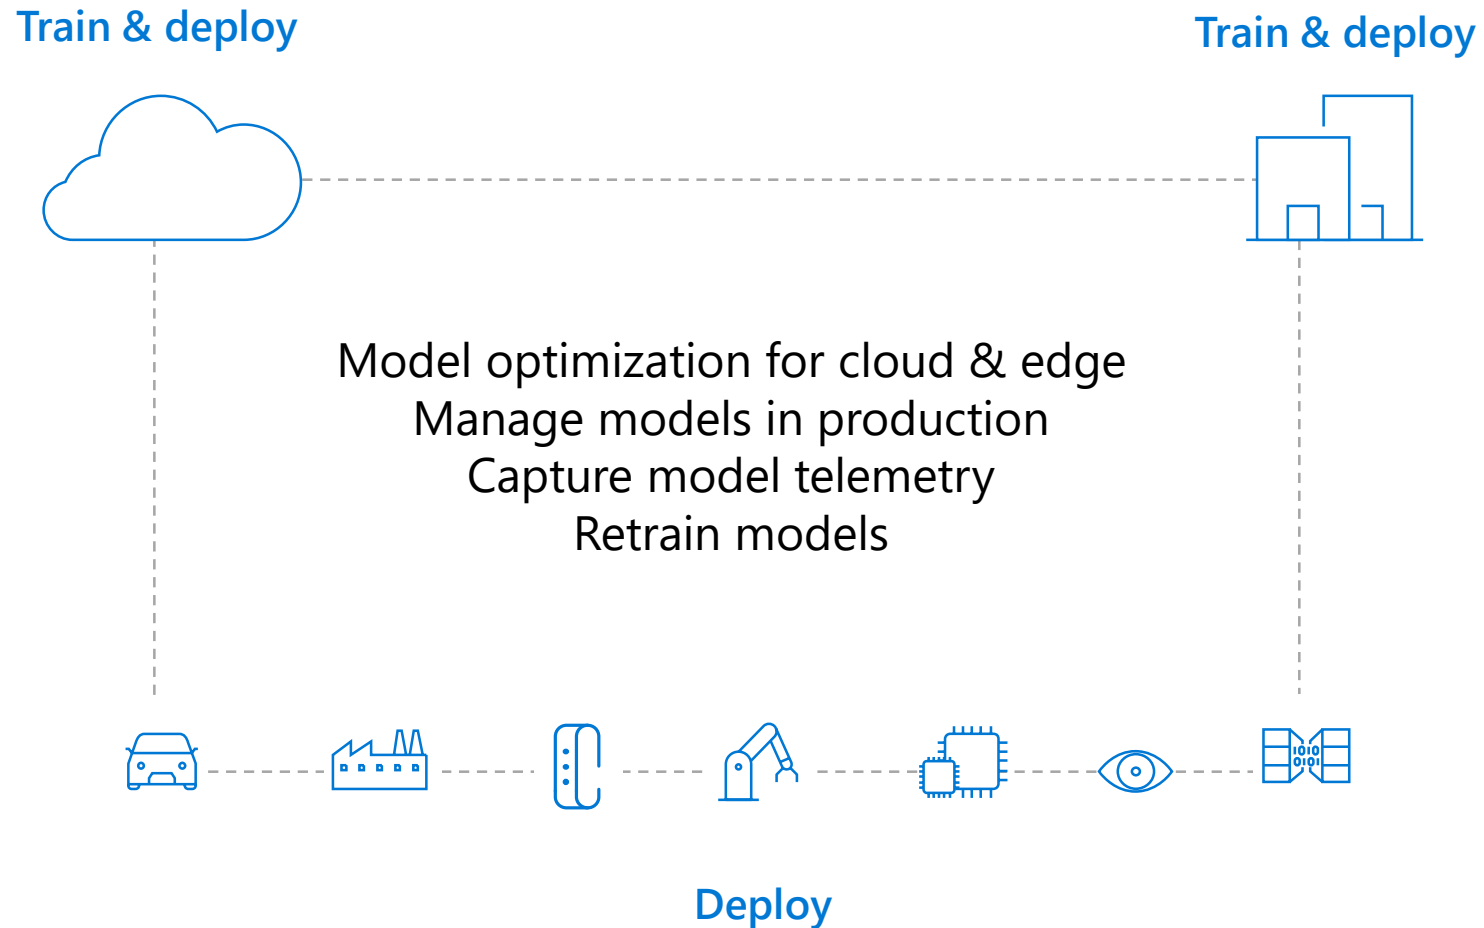
Use multiple pipelines that are reliably coordinated across heterogeneous and scalable computes and storages



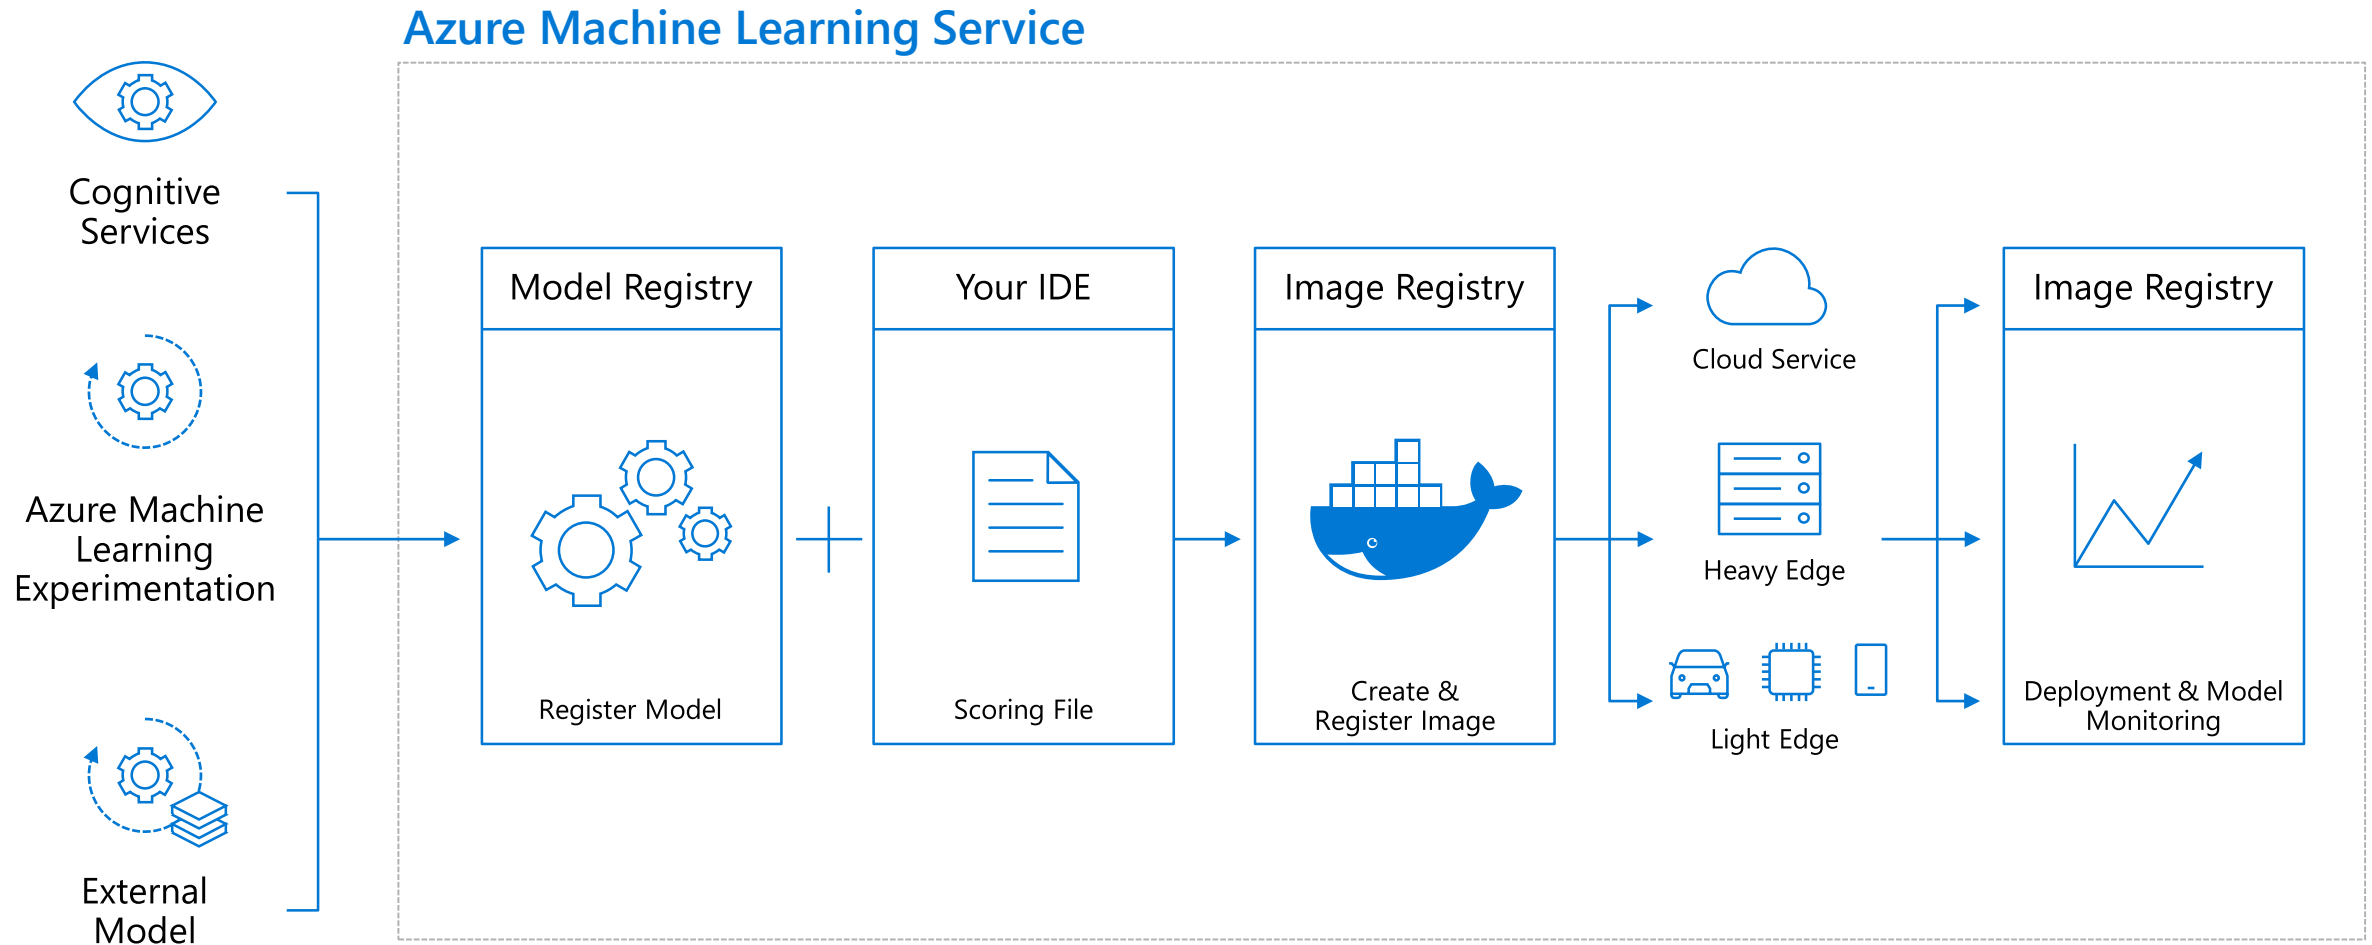
Simple deployment

Flexible deployment

Deploy and manage models on intelligent cloud and edge



Deploy Azure ML models at scale





**Support for open
source frameworks**

Popular frameworks

Use your favorite deep learning frameworks



TensorFlow



PyTorch



Scikit-Learn



MXNet



Chainer



Keras



without getting locked into one framework



ONNX

Community project created by Facebook and Microsoft

Use the best tool for the job. Train in one framework
and transfer to another for inference





Tool agnostic Python SDK

Tool Agnostic Python SDK



PyCharm



Jupyter




Visual Studio Code

Use your favorite IDEs, editors, notebooks, and frameworks



Integrate with other services like Azure Databricks




Flexibility of your local environment or curated cloud environment




Get started quickly without any complex pre-requisites

一个简单的示例

从开发工具谈起



Visual Studio Code



Azure Machine Learning

Microsoft | 271,469 | ★★★★★ | Repository | License

Visual Studio Code extension for Azure Machine Learning

Disable ▾ Uninstall

[Details](#) [Contributions](#) [Changelog](#) [Dependencies](#)

Azure Machine Learning for Visual Studio Code

Azure Machine Learning for Visual Studio Code, previously called Visual Studio Code Tools for AI**, is an extension to easily build, train, and deploy machine learning models to the cloud or the edge with [Azure Machine Learning service](#).

Getting Started

- [Installation](#)
- [Getting started with Azure Machine Learning for Visual Studio Code](#)
- [Create and manage Azure compute targets](#)
- [Train and tune models](#)
- [Deploy and manage models](#)
- [Release notes](#)

With Azure Machine Learning service, you can:


- Build and train machine learning models faster, and easily deploy to the cloud or the edge.
- Use the latest open source technologies such as [TensorFlow](#), [PyTorch](#), or [Jupyter](#).
- Experiment locally and then quickly scale up or out with large GPU-enabled clusters in the cloud.
- Speed up data science with automated machine learning and hyper-parameter tuning.
- Track your experiments, manage models, and easily deploy with integrated CI/CD tooling.

** Previous documentation and vsix installer are moved to the [archive folder](#).

Supported Operating Systems

Currently this extension supports the following 64-bit operating systems:

- Windows
- macOS



示例

Visual Studio Code interface showing a Python script named `demo1.py` and its output in the Python Interactive window.

EXPLORER

- OPEN EDITORS
 - GROUP 1
 - demo1.py
 - GROUP 2
 - Python Interactive
- AZURE DEVDAY
 - .vscode
 - gear_images
 - resized_images

demo1.py

```
1 Run Cell | Run Below
2 #%%
3 import os
4 import shutil
5 import numpy as np
6 import matplotlib.pyplot as plt
7 from PIL import Image
8
9 Run Cell | Run Above | Run Below
10 #%%
11 get_ipython().run_line_magic('matplotlib', 'inline')
12
13 Run Cell | Run Above | Run Below
14 #%%
15 imgdir = 'gear_images'
16 fig = plt.figure(figsize=(12, 16))
17 dir_num = 0
18 for root, folders, filenames in os.walk(imgdir):
19     for folder in folders:
20         # Load the first image file using the PIL library
21         file = os.listdir(os.path.join(root, folder))[0]
22         imgFile = os.path.join(root, folder, file)
23         img = Image.open(imgFile)
24         # Add the image to the figure (which will have 4 col
25         a=fig.add_subplot(4,np.ceil(len(folders)/4),dir_num
26         imgplot = plt.imshow(img)
27         # Add a caption with the folder name
28         a.set_title(folder)
29         dir_num = dir_num + 1
30
31 Run Cell | Run Above | Run Below
32 #%%
33 import os
34 import shutil
35 import numpy as np
36 import matplotlib.pyplot as plt
37 from PIL import Image
38
39 # Helper function to resize image
40 def resize_image(img, size):
41     from PIL import Image, ImageOps
42     # resize the image so the longest dimension matches our
43     img.thumbnail(size, Image.ANTIALIAS)
44     # Create a new square white background image
```

Python Interactive

Jupyter Server URI: <http://localhost:8888/?token=c58f0b05e51e451519d42ca9248a41ea99745b385a32d3b1>
Python version: 3.6.5 [Anaconda, Inc.] (default, Mar 29 2018, 13:32:41) [MSC v.1900 64 bit (AMD64)]
C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\python.exe
Jupyter Notebook Version: (5, 5, 0)

[1] import os...

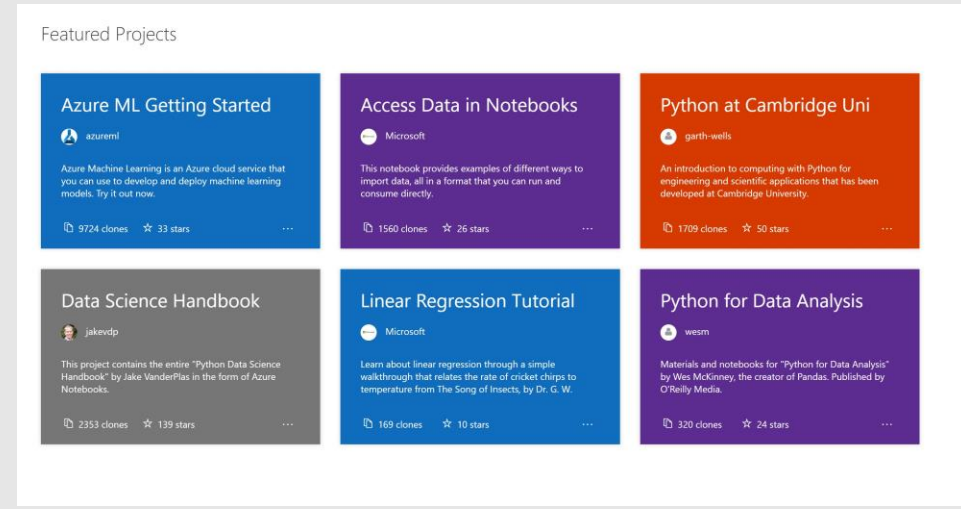
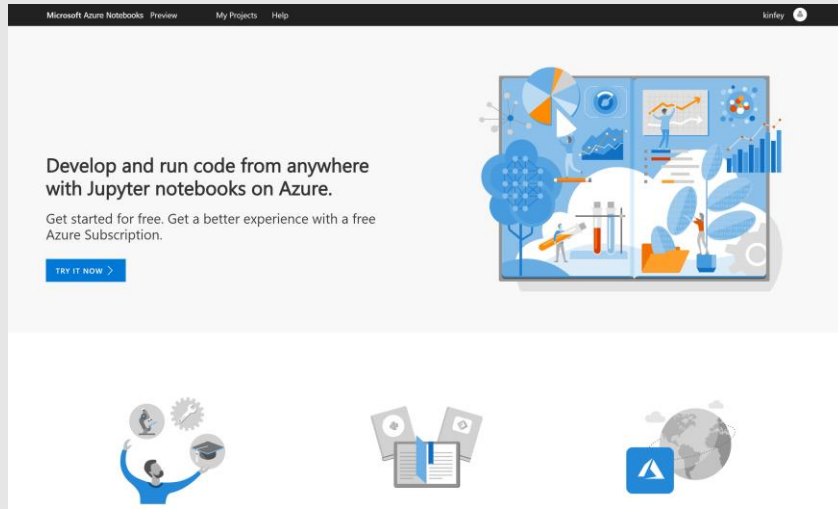
[2] get_ipython().run_line_magic('matplotlib', 'inline')

[4] imgdir = 'gear_images'...

The output displays a grid of images from the `gear_images` directory, including items like a climbing rope, a climbing shoe, a carabiner, a climbing helmet, a climbing harness, a climbing glove, a climbing jacket, a climbing helmet, a climbing harness, a climbing glove, and a climbing jacket.

Azure Notebook <https://notebooks.azure.com/>

Develop and run code from anywhere with Jupyter notebooks on Azure.



Programming Languages



Python 2

Python is a dynamically typed programming language designed by Guido Van Rossum. Much like the programming language Ruby, Python was designed to be easily read by programmers.

[LEARN ABOUT PYTHON 2 >](#)



Python 3

Python is a dynamically typed programming language designed by Guido Van Rossum. Much like the programming language Ruby, Python was designed to be easily read by programmers.

[LEARN ABOUT PYTHON 3 >](#)



R

R is an open source programming language and software environment for statistical computing and graphics that is supported by the R Foundation for Statistical Computing.

[LEARN ABOUT R >](#)



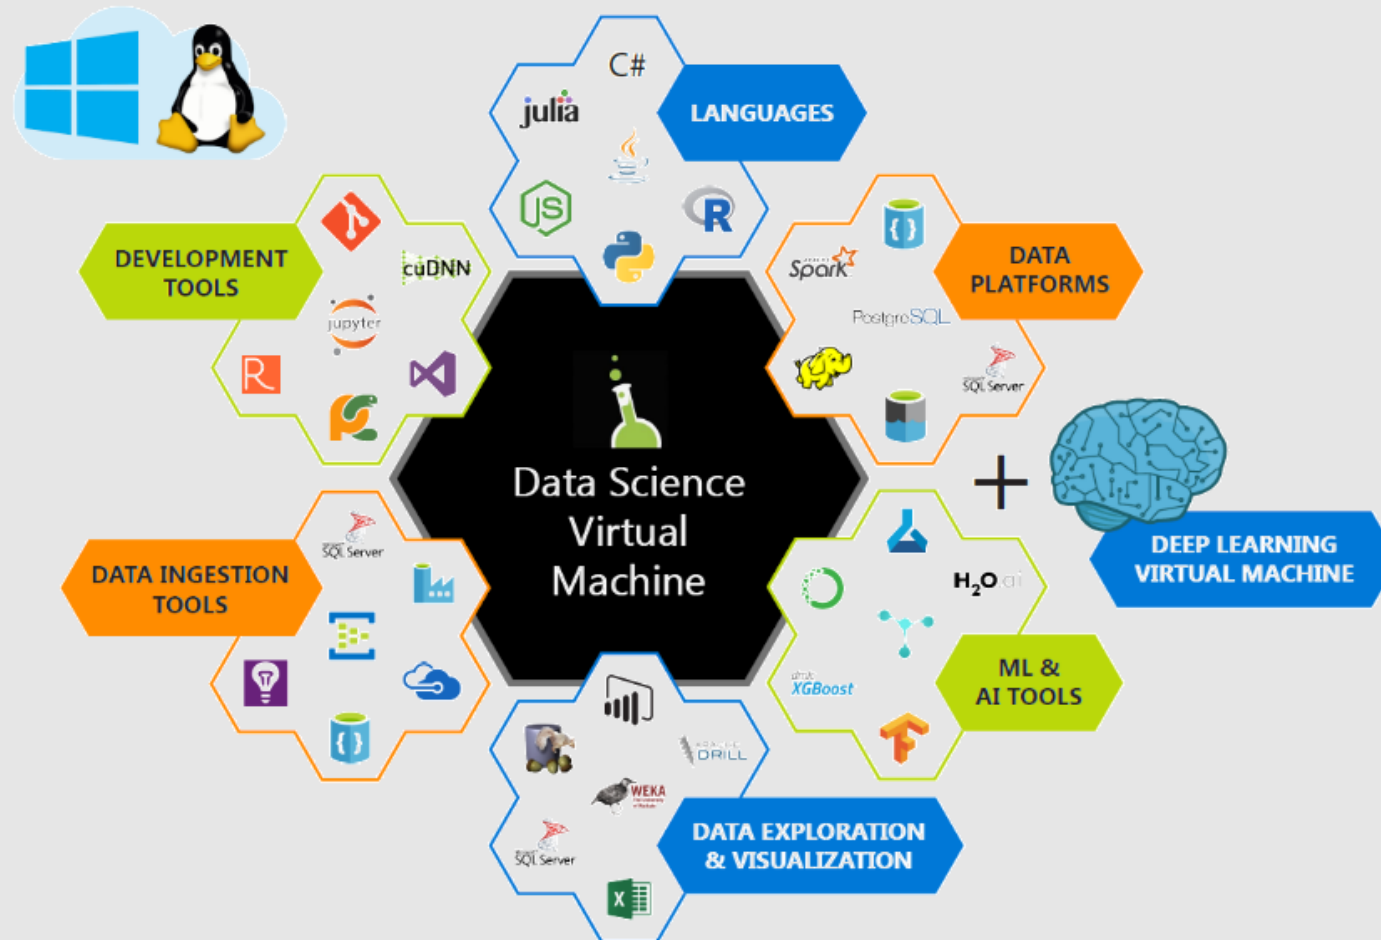
F#

F# is a strongly typed, multi-paradigm programming language that encompasses functional, imperative, and object-oriented programming methods.

[LEARN ABOUT F# >](#)

Data Science Virtual Machine(DSVM)

**Pre-Configured environments in the cloud for
Data Science & AI Modeling, Development & Deployment.**



Data Science Virtual Machine(DSVM)



DSVM – Windows Server 2016

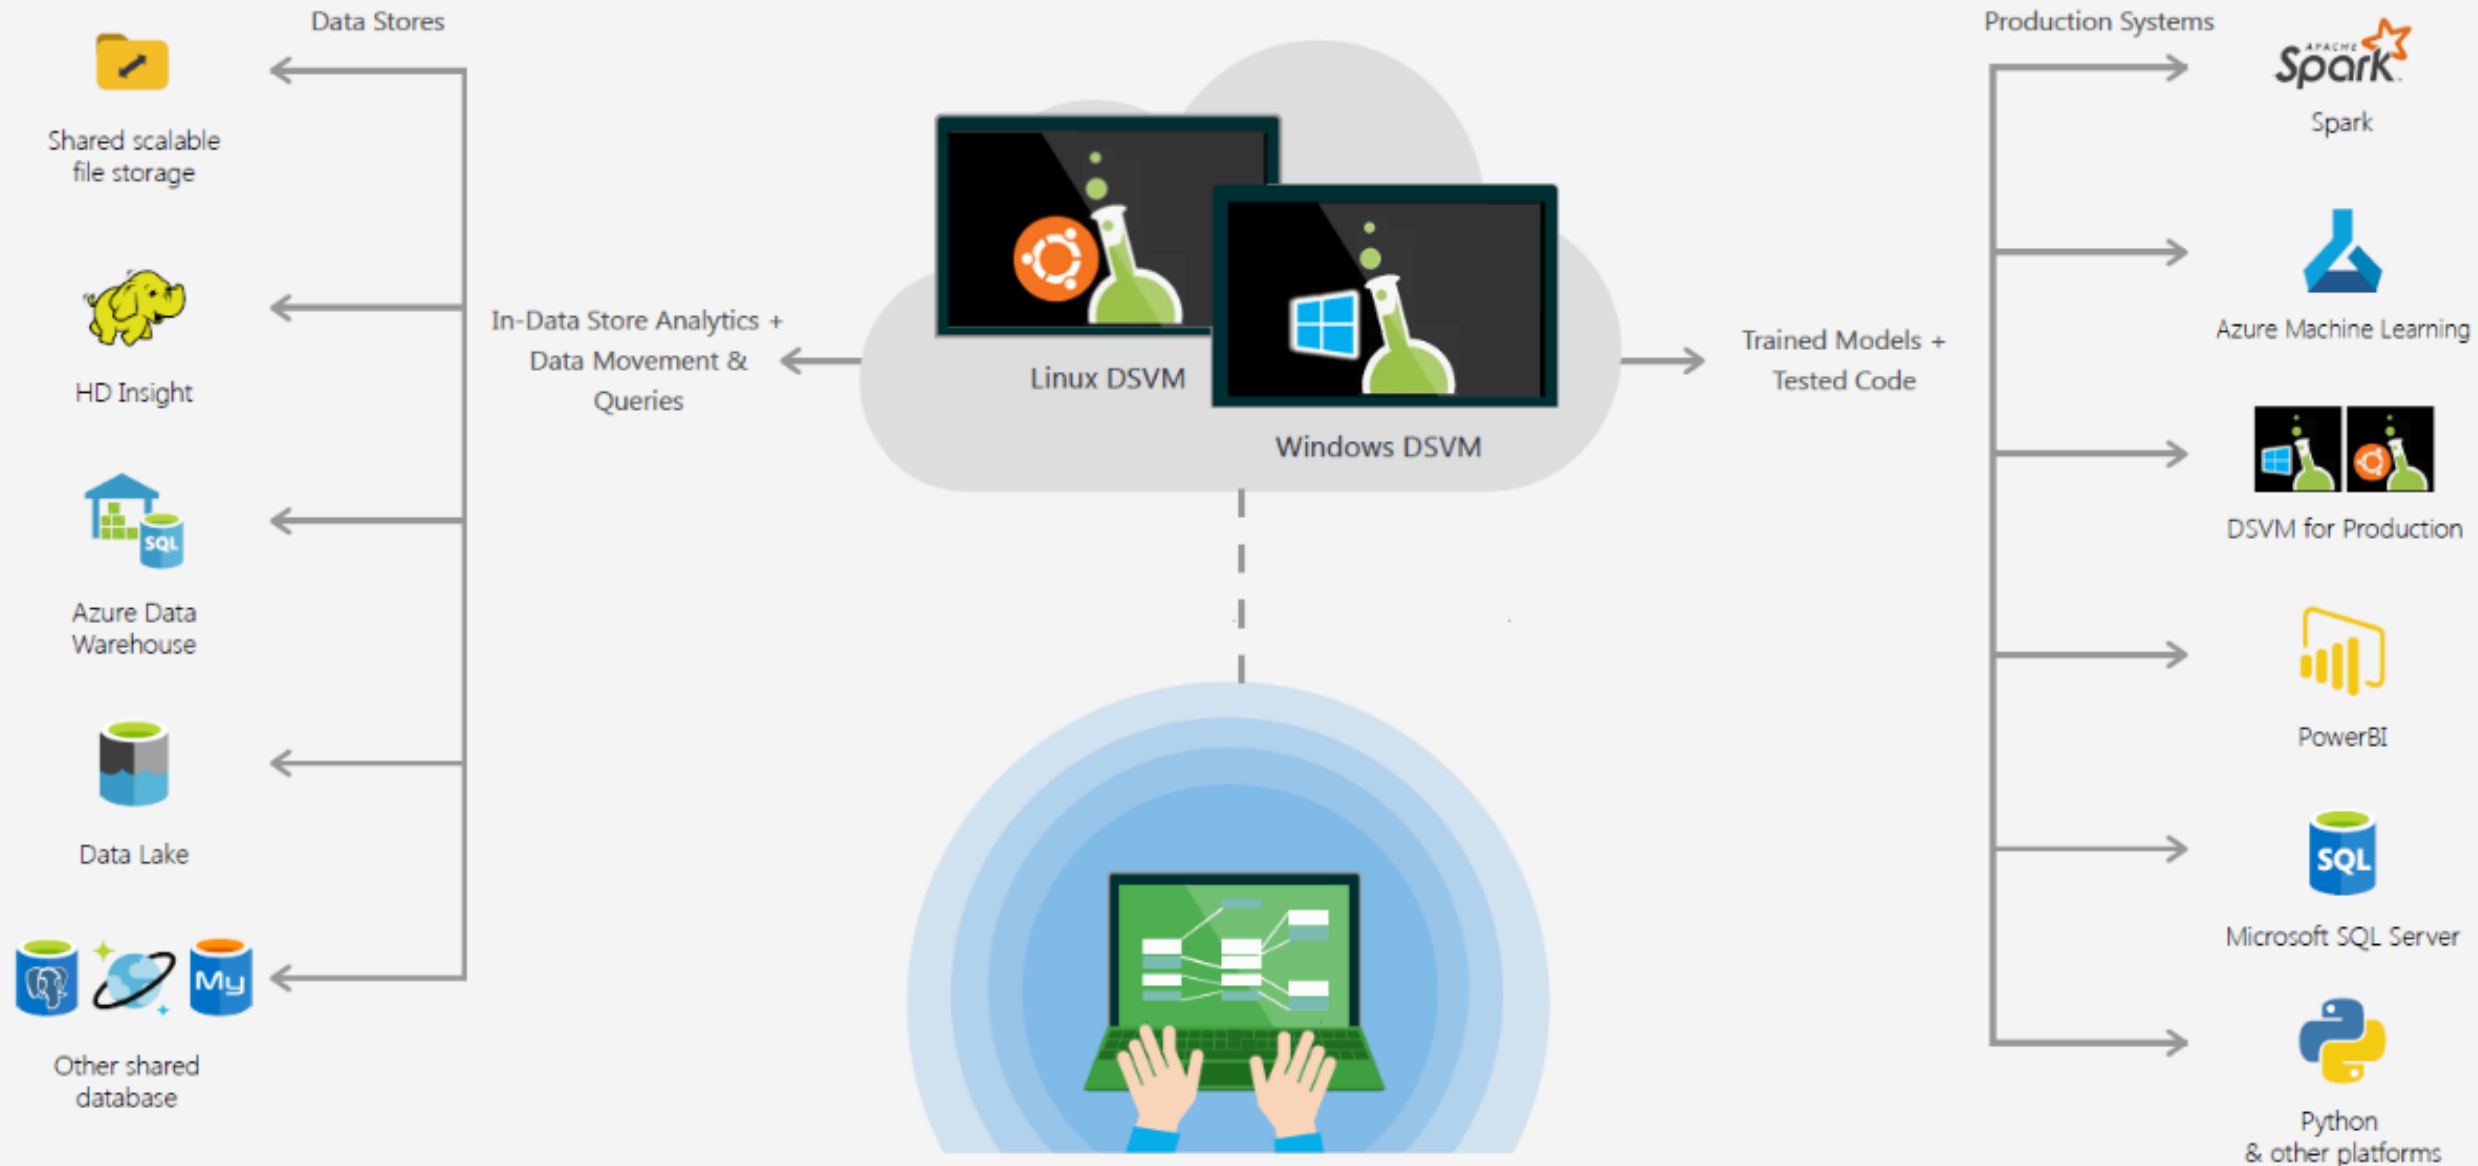


DSVM – Linux – Ubuntu



Deep Learning Virtual Machines

End-to-End AI Development Workflow using Data Science Virtual Machines (DSVM)



Demo

Jupyterhub on DSVM

jupyter lab03 Last Checkpoint: 5 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Save Add Copy Paste Undo Redo Run Stop Restart Code Edit Presentation Show Presentation

```
model.add(Dense(100, activation='softmax'))  
model.compile(optimizer=Adam, loss='categorical_crossentropy', metrics=['accuracy'])  
return model
```

In [5]:


```
import matplotlib  
matplotlib.use("Agg")  
  
from keras.preprocessing.image import ImageDataGenerator  
from keras.optimizers import Adam  
from keras.preprocessing.image import img_to_array  
from keras.utils import np_utils  
from sklearn.preprocessing import LabelBinarizer  
from sklearn.model_selection import train_test_split  
# from imagesearch.smallervggnet import SmallerVGGNet  
import matplotlib.pyplot as plt  
from imutils import paths  
import numpy as np  
import argparse  
import random  
import pickle  
import cv2  
import os
```

In [6]:

```
EPOCHS = 100  
INIT_LR = 1e-3  
BS = 32  
IMAGE_DIMS = (96, 96, 3)  
  
data = []  
labels = []  
  
print("[INFO] loading images...")  
imagePaths = sorted(list(paths.list_images('resized_images')))  
random.seed(42)  
random.shuffle(imagePaths)  
  
for imagePath in imagePaths:  
    image = cv2.imread(imagePath)  
    image = cv2.resize(image, (IMAGE_DIMS[1], IMAGE_DIMS[0]))  
    image = img_to_array(image)  
    data.append(image)
```

Azure Machine Learning service

Azure Cloud
Services



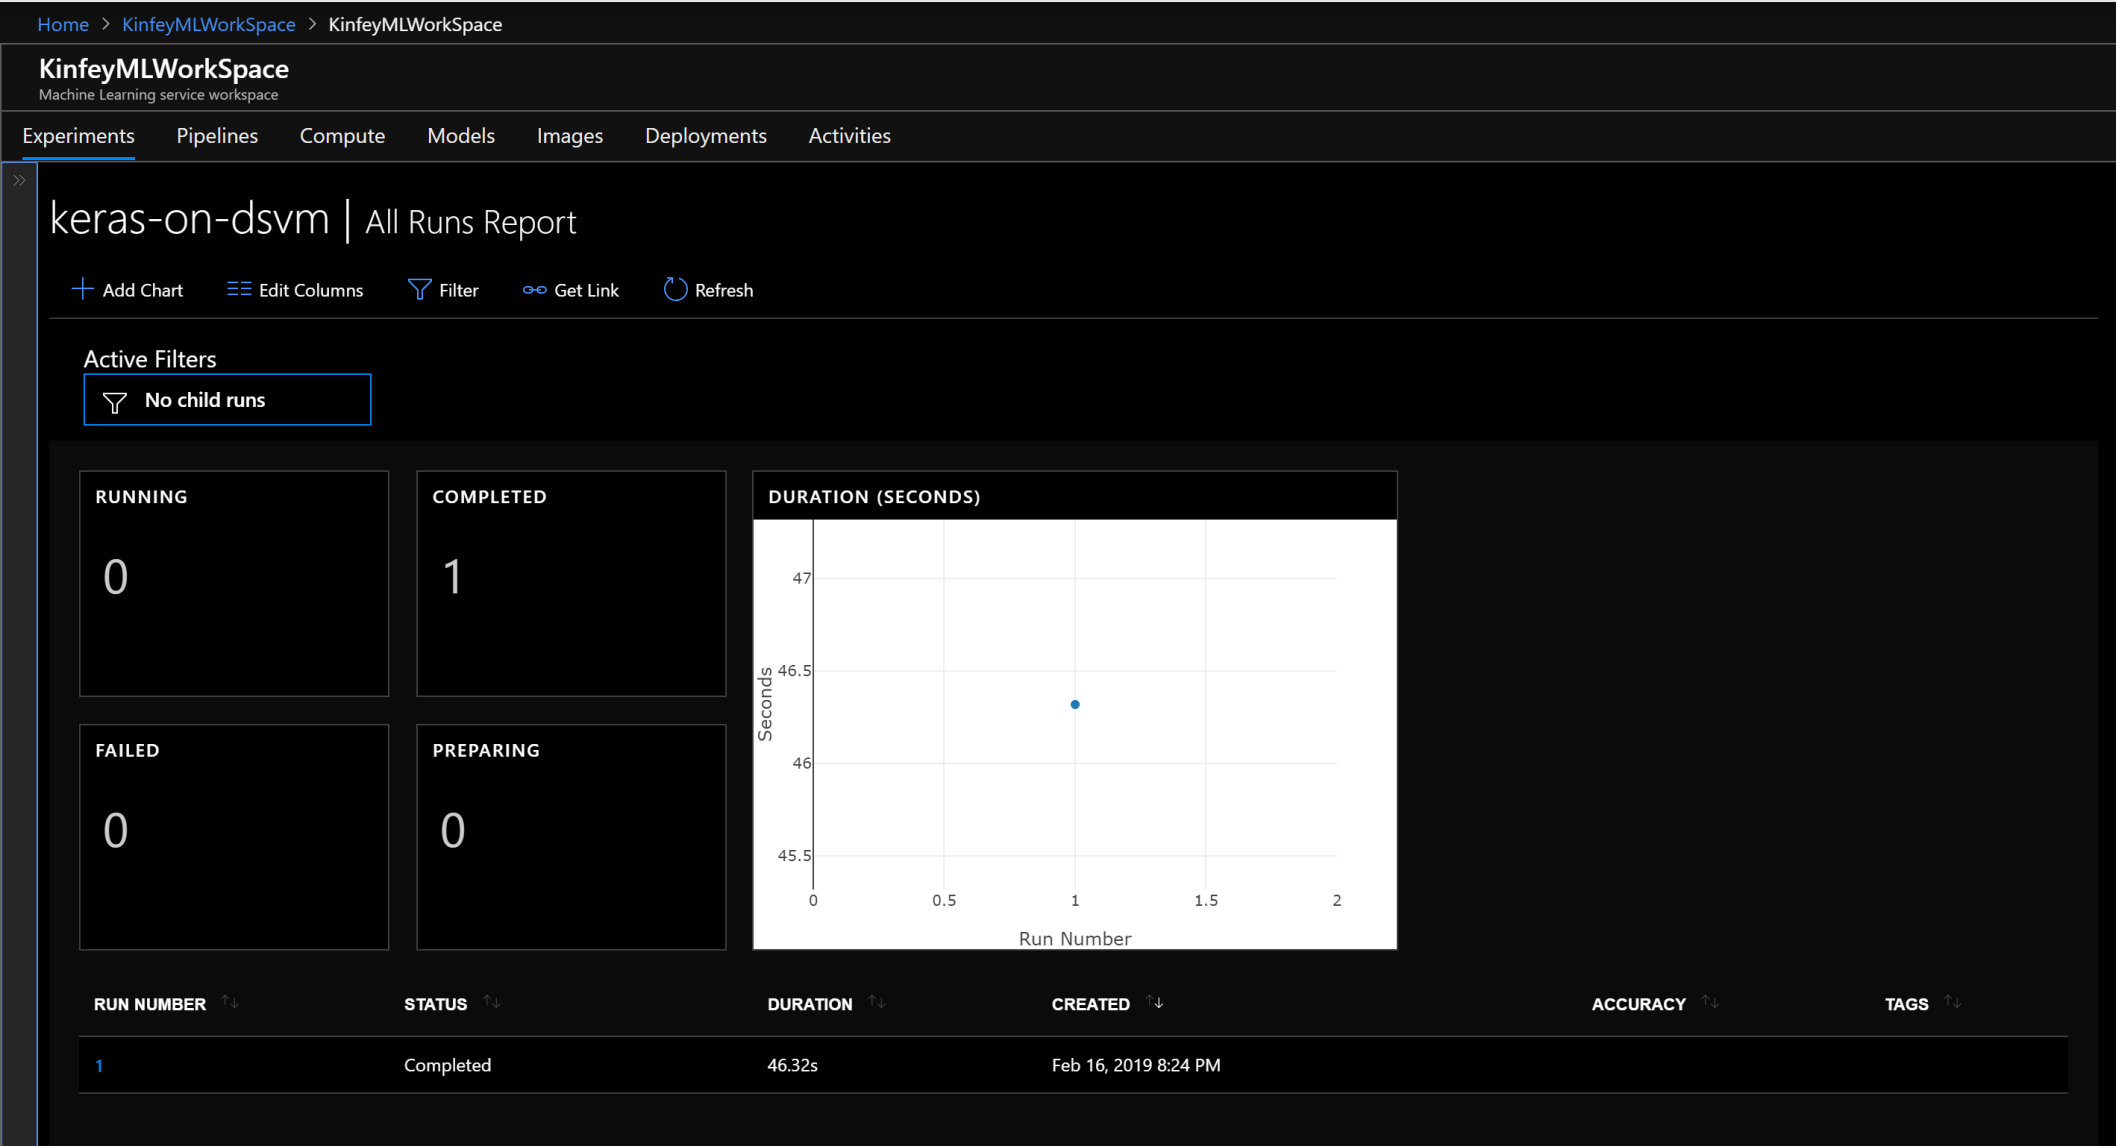
Python
SDK

帮助你完成:


- ✓ 数据准备
- ✓ 编译模型
- ✓ 训练模型

- ✓ 管理模型
- ✓ 跟踪训练
- ✓ 部署模型

Azure Machine Learning Workbench



Azure Machine Learning Tools for VS Code



Azure Machine Learning

ms-toolsai.vscode-ai Preview

Microsoft | 241,098 | ★★★★★ | Repository | License

Visual Studio Code extension for Azure Machine Learning

Disable Uninstall

Details Contributions Changelog Dependencies

Azure Machine Learning for Visual Studio Code

Azure Machine Learning for Visual Studio Code, previously called Visual Studio Code Tools for AI**, is an extension to easily build, train, and deploy machine learning models to the cloud or the edge with [Azure Machine Learning service](#).

With Azure Machine Learning service, you can:

- Build and train machine learning models faster, and easily deploy to the cloud or the edge.
- Use the latest open source technologies such as [TensorFlow](#), [PyTorch](#), or [Jupyter](#).
- Experiment locally and then quickly scale up or out with large GPU-enabled clusters in the cloud.
- Speed up data science with automated machine learning and hyper-parameter tuning.
- Track your experiments, manage models, and easily deploy with integrated CI/CD tooling.

** Previous documentation and vsix installer are moved to the archive folder.

Supported Operating Systems

Currently this extension supports the following 64-bit operating systems:

- Windows
- macOS
- Linux Ubuntu

Getting Started

- [Release notes](#)
- [Installation](#)
- [Getting started with Azure Machine Learning](#)
- [Create and manage Azure compute targets](#)
- [Train and tune models](#)
- [Deploy and manage models](#)

AZURE MACHINE LEARNING

- Windows Azure MSDN - Visual Stu...
- KinfeyMLWorkspace
 - Experiments
 - keras-on-dsvm
 - Pipelines
 - Compute
 - Models
 - Images
 - Deployments

Untitled-1 • Extension: Azure Machine Learning

```
1 Run Cell | Run All Cells
2 #%%
3 from keras import backend
4 print(backend._BACKEND)
5
6
7 Run Cell | Run All Cells
8 #%%
9 from keras.models import Sequential
10 from keras.layers import Dense
11 import numpy as np
12
13 np.random.seed(7)
14
15
16 dataset = np.loadtxt('pima-indians-diabetes.csv',delimiter=',')
17
18
19 x = dataset[:,0:8]
20 Y = dataset[:,8]
21
22 model = Sequential()
23 model.add(Dense(12, input_dim=8, activation='relu'))
24 model.add(Dense(8, activation='relu'))
25 model.add(Dense(1, activation='sigmoid'))
26
27 model.compile(loss='binary_crossentropy',optimizer='adam',metrics=['accuracy'])
28
29 model.fit(x=x,y=Y,epochs=150,batch_size=10)
30
31 scores=model.evaluate(x=x,y=Y)
32 print('\n%s : %.2f%%' % (model.metrics_names[1],scores[1]*100))
33
34
35
36 Run Cell | Run All Cells
37 #%%
38
```


1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

20

21

22

23

24

25

26

27

28

29

30

31

32

33

34

35

36

Run Cell | Run All Cells

###

from keras import backend

print(backend._BACKEND)

Run Cell | Run All Cells

###

from keras.models import Sequential

from keras.layers import Dense

import numpy as np

np.random.seed(7)

dataset = np.loadtxt('pima-indians-diabetes.csv',delimiter=',')

x = dataset[:,0:8]

Y = dataset[:,8]

model = Sequential()

model.add(Dense(12, input_dim=8, activation='relu'))

model.add(Dense(8, activation='relu'))

model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',optimizer='adam',metrics=['accuracy'])

model.fit(x=x,y=Y,epochs=150,batch_size=10)

scores=model.evaluate(x=x,y=Y)

print('\n%s : %.2f%%' % (model.metrics_names[1],scores[1]*100))

Run Cell | Run All Cells

###

Azure Machine Learning

+

🔄

Windows Azure MSDN - Visual Studio

KinfeyMLWorkSpace

Experiments

keras-on-dsvm

Pipelines

Compute

Models

Images

Deployments

Untitled-1

Extension: Azure Machine Learning

Python Interactive

Jupyter Server URI: http://localhost:8888/?token=f19ba4f243b52212b37ef26a13e374145de001acfaeb13b3

Python Version: 3.6.6 (v3.6.6:4cf1f54eb7, Jun 27 2018, 03:37:03) [MSC v.1900 64 bit (AMD64)]

C:\Program Files (x86)\Microsoft Visual Studio\Shared\Python36_64\python.exe

Jupyter Notebook Version: (5, 7, 4, '')

from keras import backend...

Using TensorFlow backend.

tensorflow

from keras.models import Sequential...

768/768 [=====] - 0s 251us/step - loss: 0.6174 - acc: 0.6992

Epoch 15/150

768/768 [=====] - 0s 242us/step - loss: 0.6015 - acc: 0.6979

Epoch 16/150

768/768 [=====] - 0s 260us/step - loss: 0.5877 - acc: 0.7018

Epoch 17/150

768/768 [=====] - 0s 264us/step - loss: 0.5850 - acc: 0.6979

Epoch 18/150

768/768 [=====] - 0s 311us/step - loss: 0.5999 - acc: 0.6901

Epoch 19/150

768/768 [=====] - 0s 259us/step - loss: 0.5806 - acc: 0.7122

Epoch 20/150

768/768 [=====] - 0s 245us/step - loss: 0.5797 - acc: 0.7240

Epoch 21/150

768/768 [=====] - 0s 260us/step - loss: 0.5712 - acc: 0.7148

Epoch 22/150

768/768 [=====] - 0s 238us/step - loss: 0.5831 - acc: 0.6992

Epoch 23/150

768/768 [=====] - 0s 276us/step - loss: 0.5751 - acc: 0.7122

Epoch 24/150

768/768 [=====] - 0s 299us/step - loss: 0.5704 - acc: 0.7266

Epoch 25/150

768/768 [=====] - 0s 255us/step - loss: 0.5577 - acc: 0.7383

Epoch 26/150

768/768 [=====] - 0s 239us/step - loss: 0.5720 - acc: 0.7044

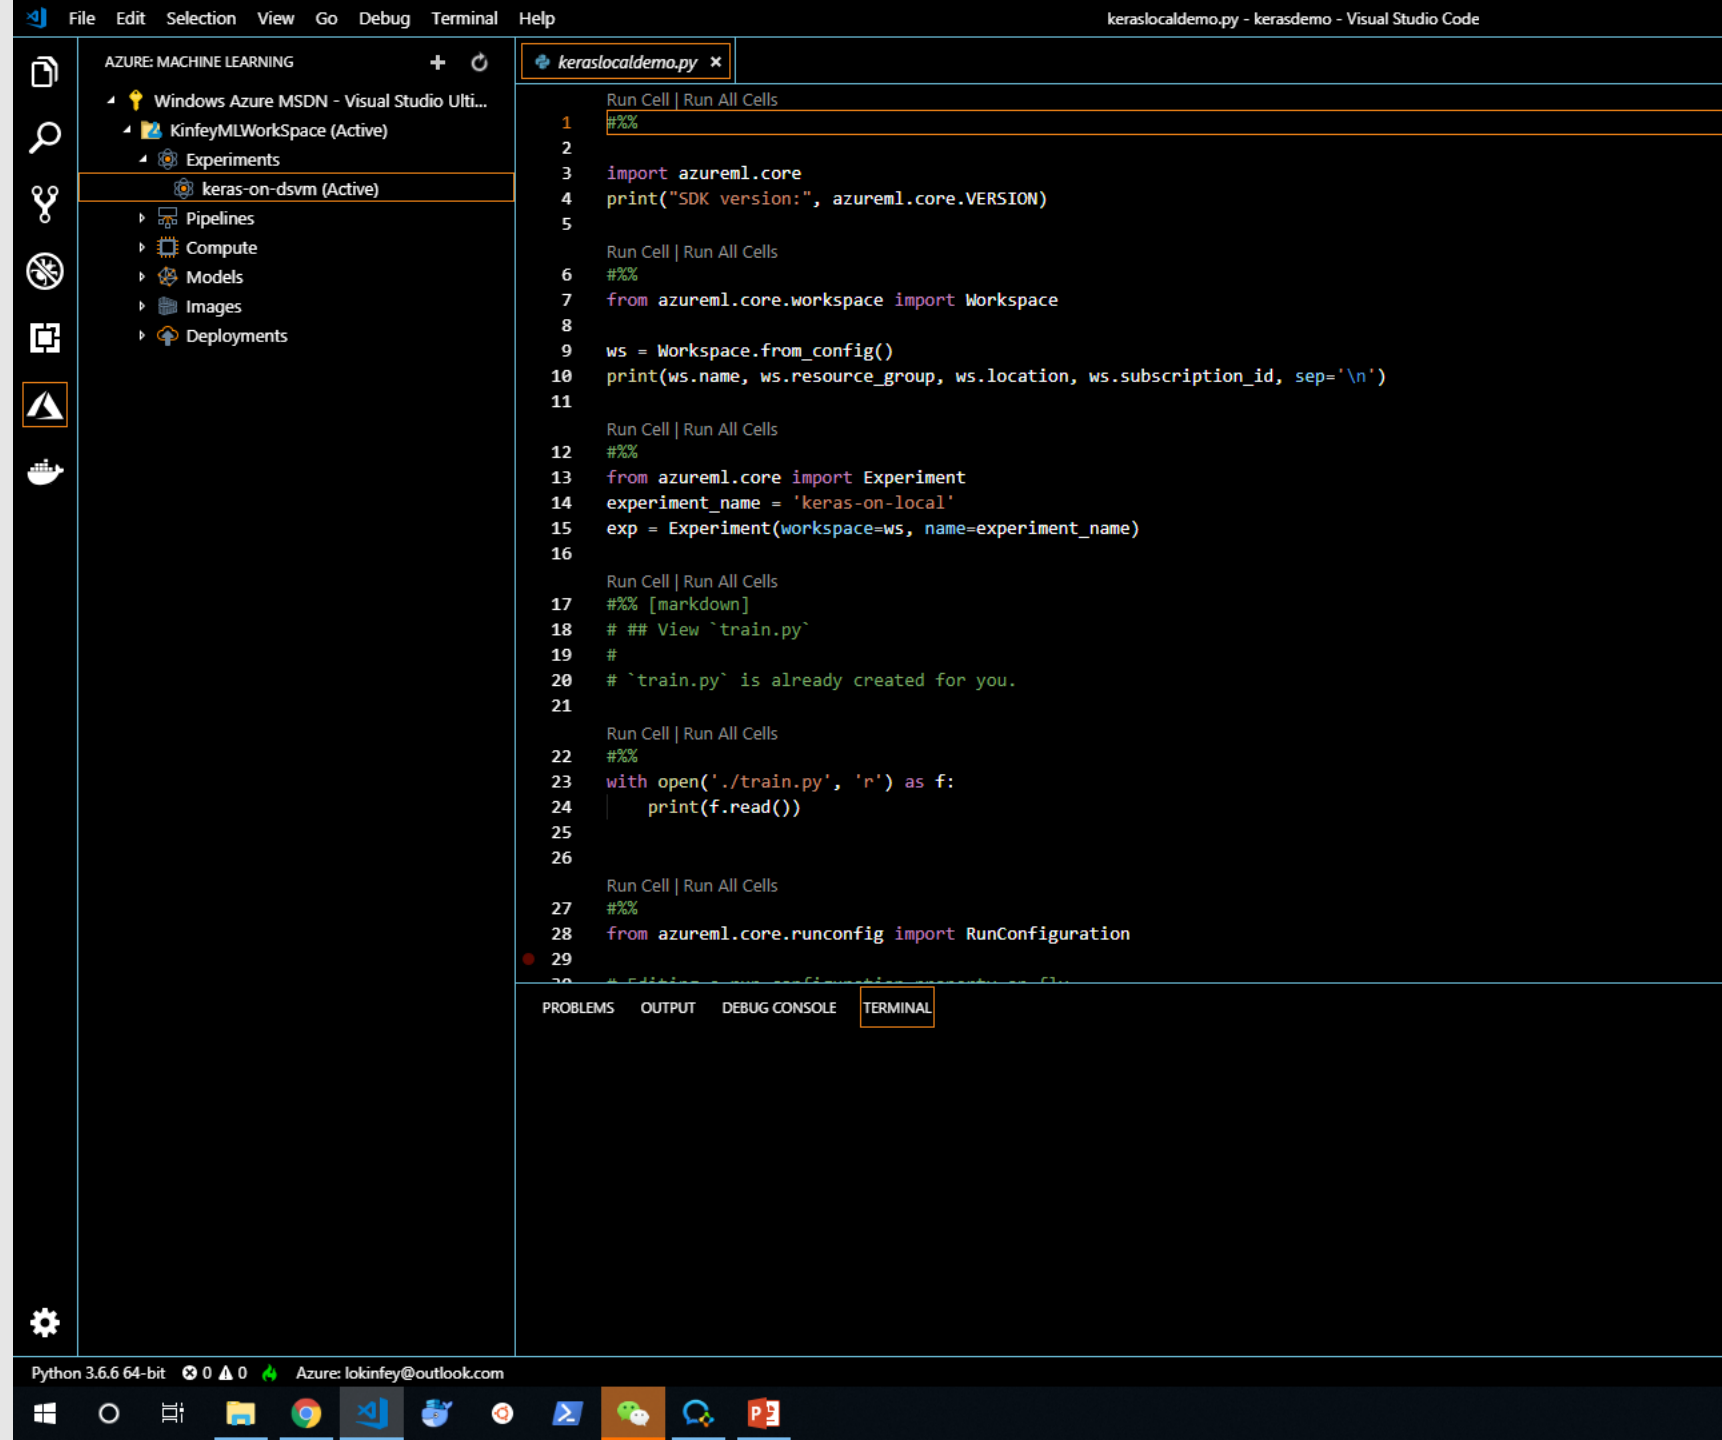
Epoch 27/150

768/768 [=====] - 0s 216us/step - loss: 0.5554 - acc: 0.7500

>>> Shift-enter to run

Demo

Azure ML Workbench Demo



特别感谢

